



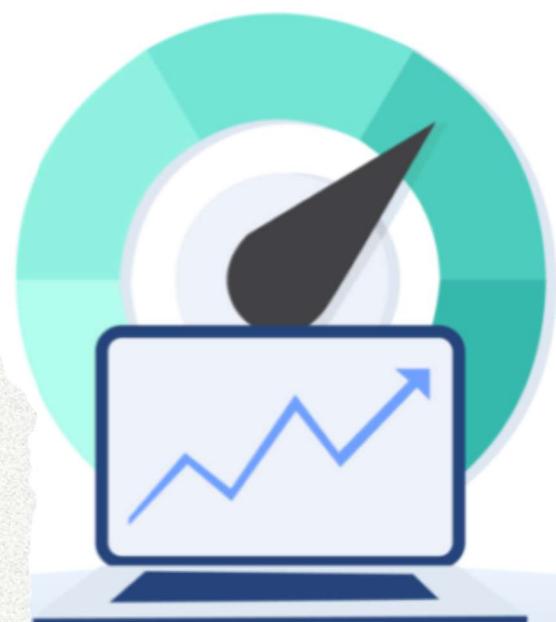
PYTHON PERFORMANCE TUNING AND OPTIMIZATION

MODULE 11

1



APPLICATION PERFORMANCE



2

What Is Performance?

- > What are some different aspects of performance?
 - Memory footprint
 - Startup time
 - Scalability
 - Responsiveness
 - Throughput
 - Efficiency
 - Utilization
 - Latency
 - Capacity
 - Degradation

3

Memory Footprint

- > The memory footprint is the amount of memory used by your application and your JVM on a system. Considerations include the following:
 - Does the application run on a dedicated system?
 - Does the application share the machine with other applications and/or JVMs?
 - As memory footprint grows, does it affect Virtual Memory?
 - Accessing data in virtual memory is much slower than accessing data in RAM.

4

Startup Time

- > *Startup time*: The time taken for an application to start
 - For client applications, this can be very important.
 - For server applications, it is less important.
 - Time ‘til Performance



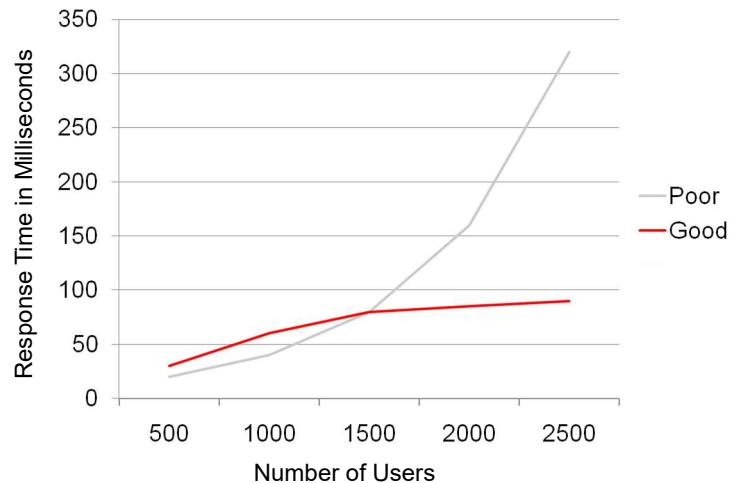
5

Scalability

- > *Scalability*: How well an application performs as the load on it increases
 - An application may perform well in development but poorly after it is deployed.
 - If an application’s response times grow *exponentially* when under load, its scalability is poor.
 - If an application’s response times grow *linearly* when under load, its scalability is good.
 - Scalability can be measured at a number of different levels in a complex system.

6

Application Scalability



7

Responsiveness

- > **Responsiveness:** How quickly an application or system responds with a requested piece of data
 - Examples
 - How quickly a desktop UI responds to an event
 - How fast a website returns a page
 - How fast a database query is returned
 - Large pause times are not acceptable.
 - The focus is on responding in short periods of time.

8

Throughput

- > *Throughput*: Focusing on maximizing the amount of work by an application in a specific period of time
 - Examples
 - The number of instructions per second that can be executed
 - The number of jobs that a batch program can complete in an hour
 - The number of database queries that can be completed in an hour
 - High pause times are acceptable.
 - The focus is on how much work can be done over longer periods of time.

9

Low Latency

Data passing	Latency	Light over a fibre	Throughput on at a time
Method call	Inlined: 0 Real call: 50 ns.	10 meters	20,000,000/sec
Shared memory	200 ns	40 meters	5,000,000/sec
SYSV Shared memory	2 μ s	400 meters	500,000/sec
Low latency network	8 μ s	1.6 km	125,000/sec
Typical LAN network	30 μ s	6 km	30,000/sec
Typical data grid system	800 μ s	160 km	1,250/sec
60 Hz power flickers	8 ms	1600 km	120/sec
4G request latency in UK	55 ms	11,000 km	18/sec

10

What is an ultra low GC?

- > Another reason to use ultra low garbage rates is your caches are not being filled with garbage and they work much more efficiently and consistently.
- > If you have a web server which is producing 300 MB/s of garbage, this means less than 5% of the time you will be pausing for a GC.
- > However, it does mean that you could be filling a 32 KB L1 CPU in around 0.1 milli-seconds. Your L2 cache fills with garbage every milli-second.

11

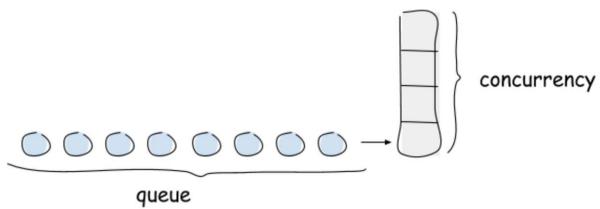
So, without a GC pause, no more pauses?

- > A high percentage of pauses are not from the GC.
- > The biggest ones are, but take away GC pauses and still see IO delays. These can be larger than GC pauses.
 - Network delays
 - Waiting for databases
 - Disk reads/writes
- > OS interrupts
 - It is not uncommon for your OS to stop your process for 5 ms or more.
- > Lock contention pauses.

12

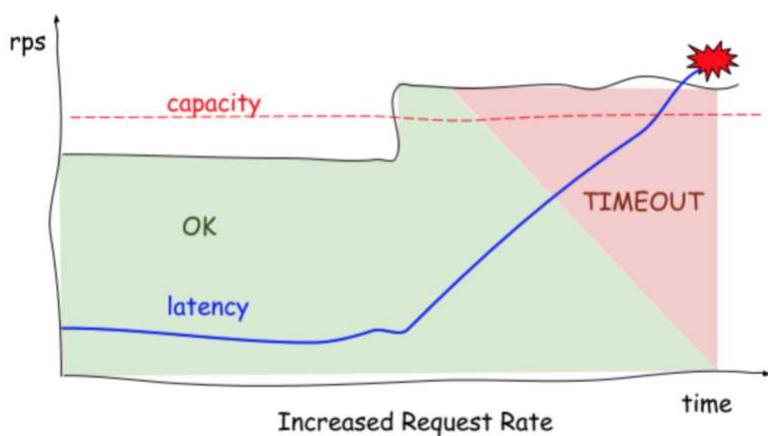
Optimizing for latency instead of throughput

- > Average latency is largely an inverse of your throughput and no better.
- > Using standard deviation for average latency is misleading at best as the distribution for latencies are not a normal distribution or anything like it.
- > From Little's Law
 - Average latency = concurrency / throughput



13

Optimizing for latency instead of throughput



14

Capacity

- > The capacity is the amount of work parallelism a system possesses
 - The number of units of work (e.g., transactions) that can be simultaneously ongoing in the system.
- > Capacity is obviously related to throughput, and we should expect that as the concurrent load on a system increases, throughput (and latency) will be affected.
- > For this reason, capacity is usually quoted as the processing available at a given value of latency or throughput.

15

Utilization

- > One of the most common performance analysis tasks is to achieve efficient use of a system's resources.
- > Ideally, CPUs should be used for handling units of work, rather than being idle (or spending time handling OS or other housekeeping tasks).
- > Depending on the workload, there can be a huge difference between the utilization levels of different resources.
- > For example, a computation-intensive workload (such as graphics processing or encryption) may be running at close to 100% CPU but only be using a small percentage of available memory.

16

Efficiency

- > Dividing the throughput of a system by the utilized resources gives a measure of the overall efficiency of the system.
- > Intuitively, this makes sense, as requiring **more resources** to produce **the same throughput** is one useful definition of being **less efficient**.
- > It is also possible, when one is dealing with larger systems, to use a form of **cost accounting to measure efficiency**.
- > If solution **A** has a total dollar cost of ownership (TCO) twice that of solution **B** for the same throughput then it is, clearly, **half as efficient**.

17

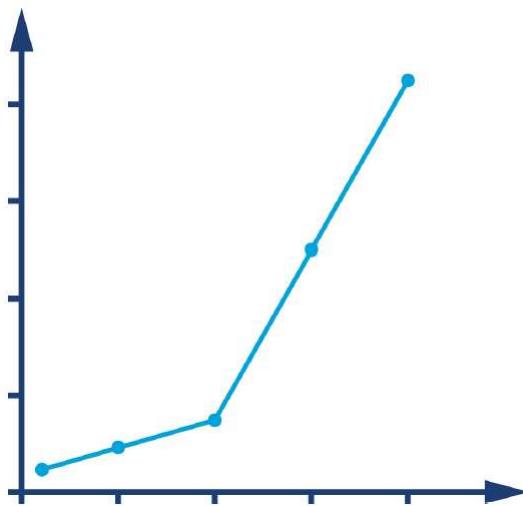
Degradation

- > If we increase the load on a system, either by increasing the number of requests (or clients) or by increasing the speed requests arrive at, then we may see a change in the observed latency and/or throughput.
- > Note that this change is dependent on utilization.
- > If the system is underutilized, then there should be some slack before observables change, but if resources are fully utilized then we would expect to see throughput stop increasing, or latency increase.
- > These changes are usually called the degradation of the system under additional load.

18

Reading Performance Graphs

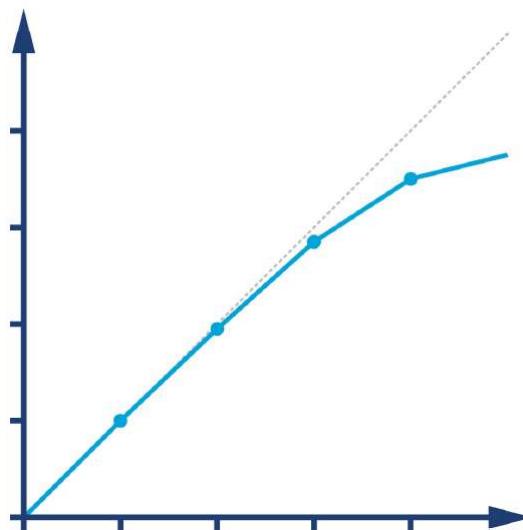
- > Latency under increasing load



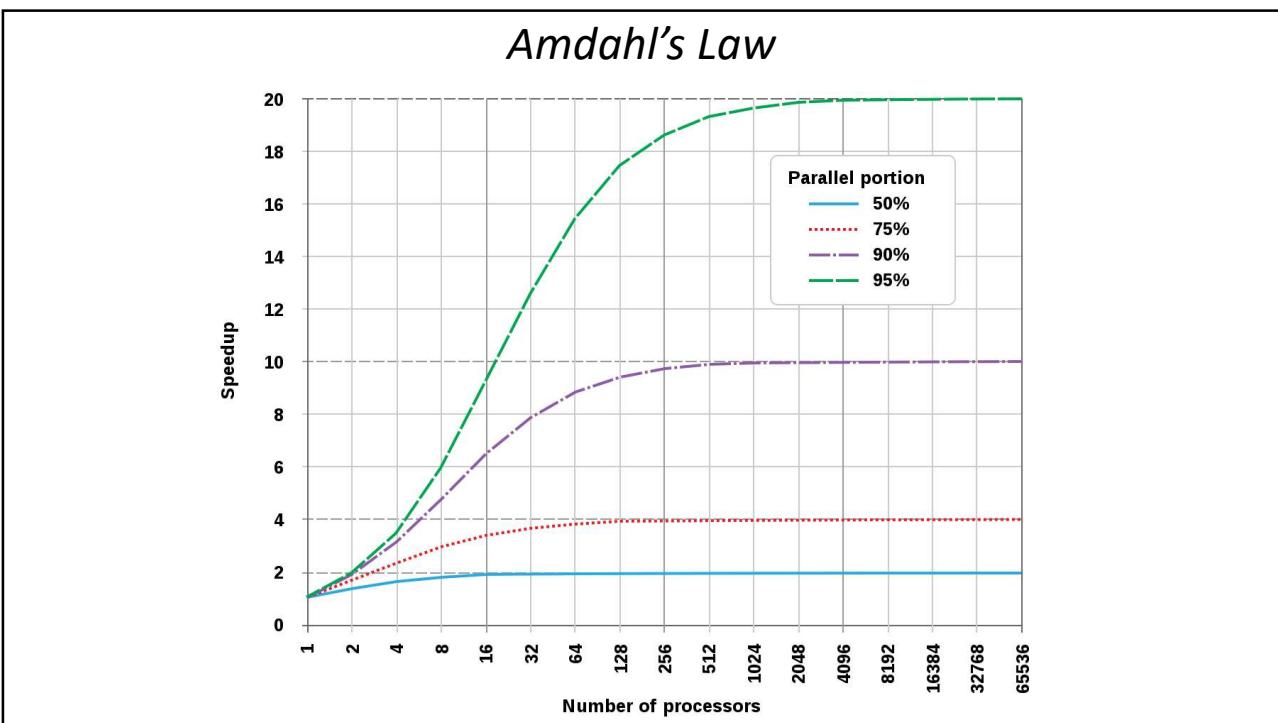
19

Reading Performance Graphs

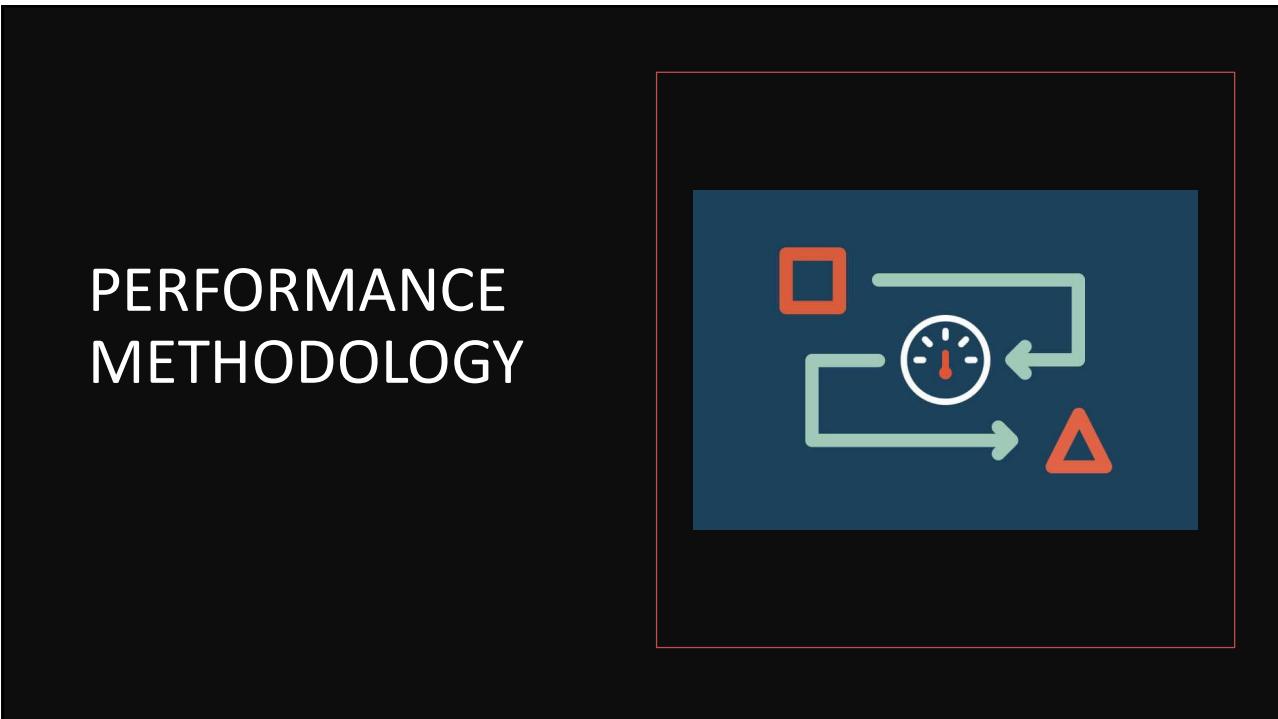
- > Latency under increasing load: *Near-linear scaling*



20



21



22

Performance Methodology

- > Performance methodology
 - Monitor
 - Profile
 - Tune
- > Definitions
- > Incorporating performance methodology into development



23

Performance Monitoring

- > Definition: An act of non-intrusively collecting or observing performance data from an operating or running application.
 - For troubleshooting
 - Identify problems
 - Determine problem characteristics
 - For development
 - Test application for meeting requirements
 - Responsiveness or throughput
 - Operating system and JVM tools



24

Performance Profiling

- > Definition: An act of collecting or observing performance data from an operating or running application.
 - Usually more intrusive than monitoring
 - Usually a narrower focus than monitoring
 - Commonly done in qualification, testing, or development environments.



25

Performance Tuning

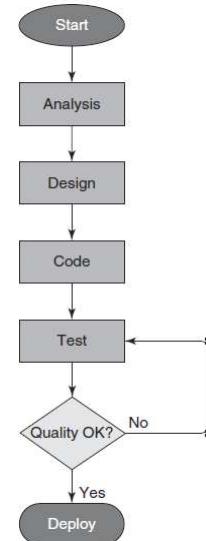
- > Definition: An act of changing tunables, source code and/or configuration attributes for the purposes of improving application responsiveness and/or application throughput.
 - Relies on performance requirements
 - Requires monitoring and/or profiling activities
 - Performance goals
 - Responsiveness
 - Throughput



26

Forces

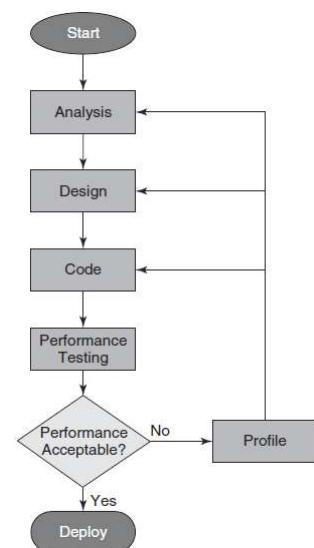
- > The traditional software development process consists of four major phases: analysis, design, coding, and testing



27

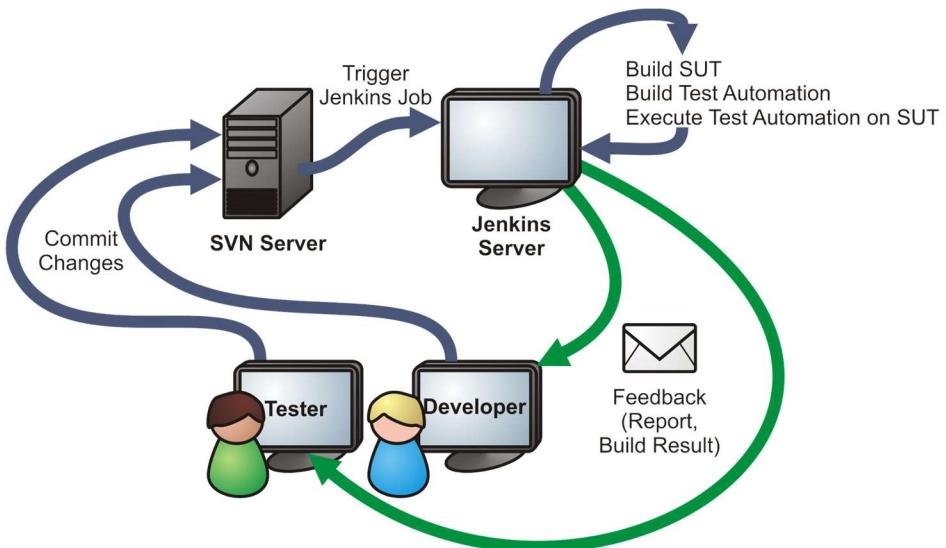
Wilson & Kesselman's Performance Process

- > Integrate performance testing into automated build processes
- > Performance regression
- > Track performance at each coding change committed to the source code base



28

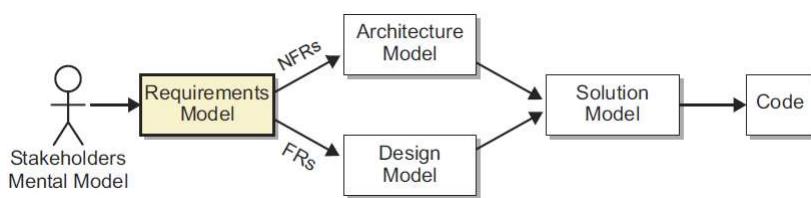
Continuous Performance Validation



29

Systemic-Quality-Driven

- > Systemic qualities are requirements on the system that are non-functional or related to the quality of service.
- > Examples include:
 - Performance – Such as responsiveness and latency
 - Reliability – The mitigation of component failure
 - Scalability – The ability to support additional load, such as more users
- > Systemic qualities drive the architecture of the software.



30

NFRs

- > What is the expected throughput of the application?
- > What is the expected latency between a stimulus and a response to that stimulus?
- > How many concurrent users or concurrent tasks shall the application support?
- > What is the accepted throughput and latency at the maximum number of concurrent users or concurrent tasks?
- > What is the maximum worst case latency?
- > What is the frequency of garbage collection induced latencies that will be tolerated?

31

Two Approaches: Top Down and Bottom Up

- > There are two commonly accepted approaches to performance analysis:

1. Top down

- Focuses on the top level of the application
- Drills down the software stack looking for problem areas and optimization opportunities

2. Bottom up

- Begins at the lowest level of the software stack
- At the CPU level looking at statistics such as CPU cache misses, inefficient use of CPU instructions
- What constructs or idioms are used by the application

32

Top-Down Approach

- > The most common approach utilized for performance tuning
- > Commonly used when you have the ability to change the code at the highest level of the application software stack
- > You begin by monitoring the application of interest under a load at which a stakeholder observes a performance issue
- > An application is continuously monitored and as a result of a change in the **application's configuration/typical load** the application experiences a degradation in performance
- > **Performance** and **scalability** requirements for the application change and the application in its current state cannot meet those new requirements.

33

Top-Down Approach

- > Monitoring the application while it is running under a load of particular interest is the first step in a top-down approach
- > Monitoring activity may include
 - Observing operating system level statistics,
 - Python interpreter statistics,
 - Application performance instrumentation statistics.

34

Top-Down Approach

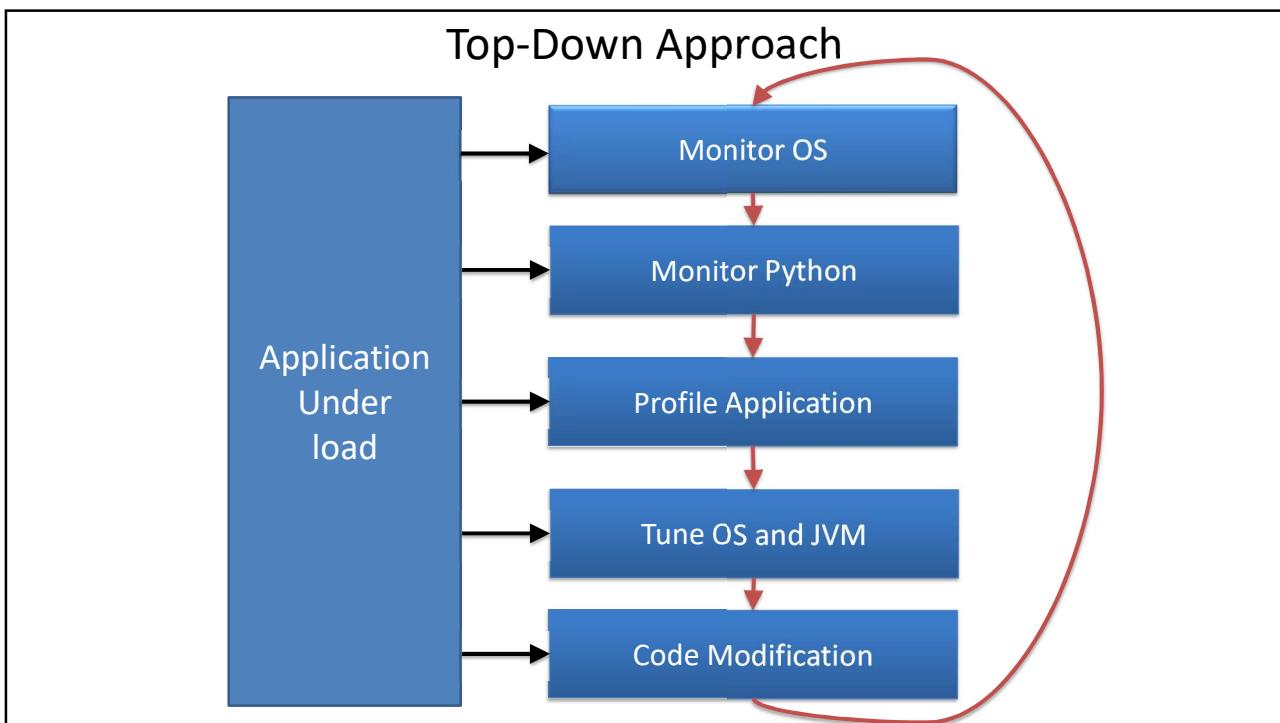
- > Tools and Techniques
 - Monitoring tools: Grafana, Prometheus.
 - Profilers: py-spy, cProfile, line_profiler.
- > Metrics to Monitor
 - Response time, throughput, memory usage, CPU utilization.

35

Top-Down Approach

- > Code Profiling
 - Use profilers to understand where time is spent.
 - Example: cProfile output for function-level performance.
- > Log Analysis
 - Review application logs for delays.
- > Database Query Analysis
 - Check query execution plans and latencies.

36

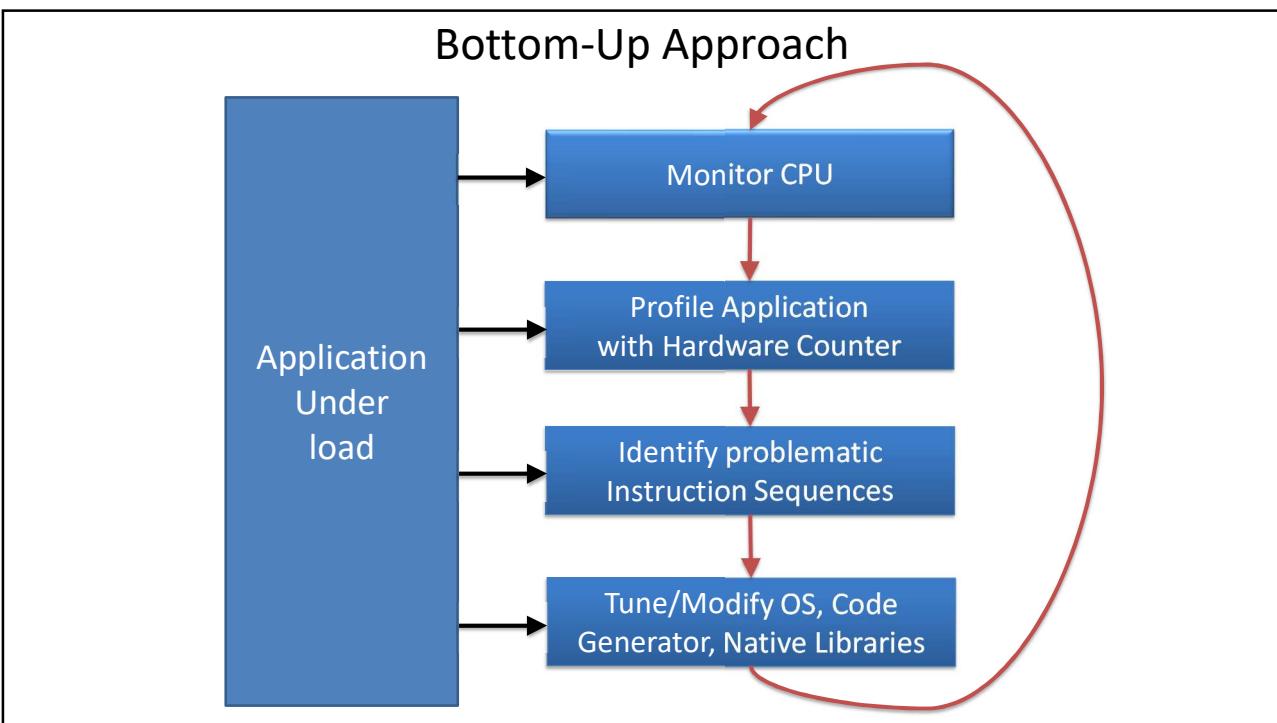


37

Bottom-Up Approach

- > The bottom-up approach is most used by performance specialists to *improve* the performance of an application on *one platform* relative to *another*
 - where differences exists in the underlying
 - CPU,
 - CPU architecture,
 - Number of CPUs.

38



39

Choosing Right Platform & Evaluating a System

> *Choosing the Right CPU Architecture*

ARM, MIPS, x86, RISC-V

ASIC, SoC, ISA

GPU, TPU

> *Evaluating a System's Performance*

40

OPERATION SYSTEM PERFORMANCE MONITORING



41

Objectives

After completing this lesson, you should be able to:

- > Monitor CPU usage
- > Monitor network I/O
- > Monitor disk I/O
- > Monitor virtual memory usage
- > Monitor processes including lock contention

42

Why Are We Monitoring?

- > Get a sense of what the problem is
- > Identify the poor performance symptoms
- > Based on the symptoms, diagnose the problem

43

Modern Hardware

- > Memory
- > Cache
- > Translation Lookaside Buffer
- > Branch Prediction and Speculative Execution
- > Hardware Memory Models

44

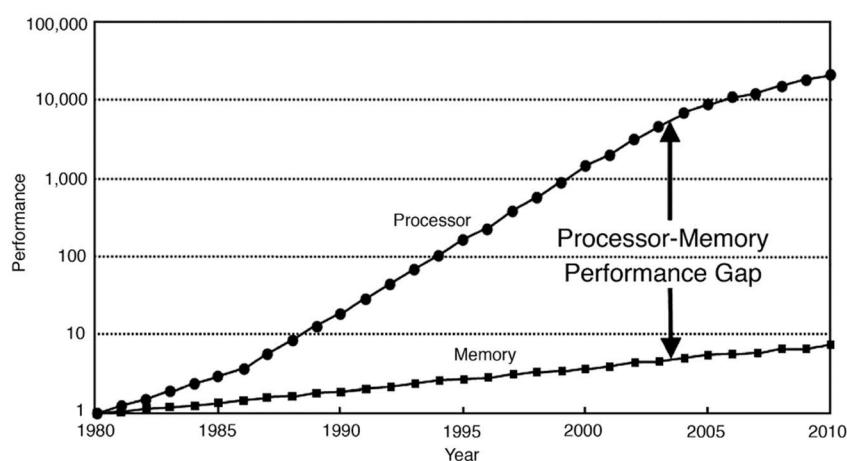
Memory

- > Moore's Law: *The number of transistors on a mass-produced chip roughly doubles every 18 months.*
- > The exponentially increasing number of transistors was initially used for faster and faster clock speed
- > Faster clock speed means more instructions completed per second
- > Faster chips require a faster stream of data to act upon
- > Over time main memory could not keep up with the demands of the processor core for fresh data

45

Memory

- > if the CPU is waiting for data, then faster cycles don't help, as CPU will just have to idle until the required data arrives.



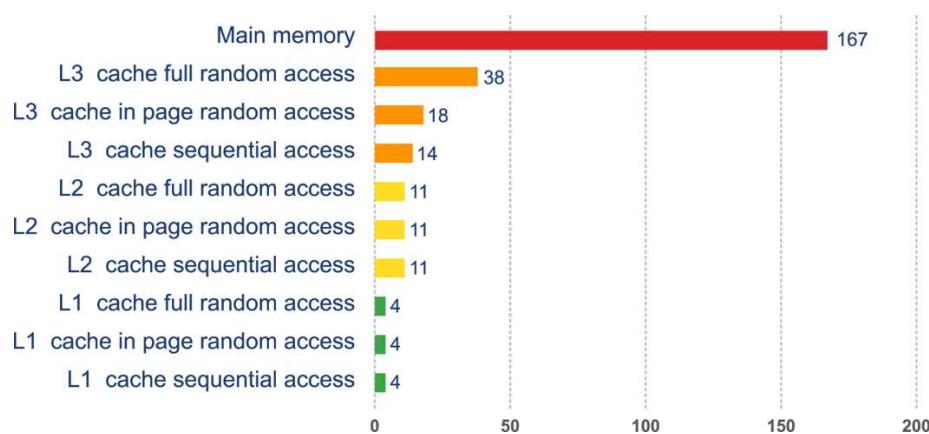
46

Memory Caches

- > To solve this problem, CPU caches were introduced.
- > These are memory areas on the CPU that are slower than CPU registers, but faster than main memory.
- > The idea is for the CPU to fill the cache with copies of often-accessed memory locations rather than constantly having to re-address main memory

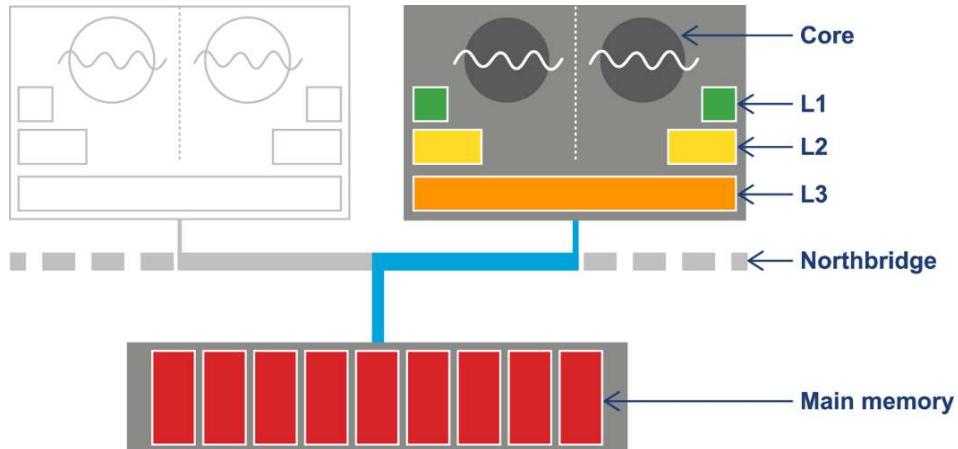
47

Memory Caches



48

Memory Caches



49

Memory Caches

- > Although the addition of a caching architecture hugely improves processor throughput, it introduces a new set of problems.
- > These problems include determining how memory is fetched into and written back from the cache.
- > The solutions to this problem are usually referred to as ***cache consistency protocols***.

50

Cache Consistency Protocols

- > At the lowest level, a protocol called **MESI** (and its variants) is commonly found on a wide range of processors.
- > It defines four states for any line in a cache.
- > Each line (usually 64 bytes) is either:
 - **Modified** (but not yet flushed to main memory)
 - **Exclusive** (present only in this cache, but does match main memory)
 - **Shared** (may also be present in other caches; matches main memory)
 - **Invalid** (may not be used; will be dropped as soon as practical)

51

Cache Consistency Protocols

- > *MESI allowable states between processors*

M E S I

M - - - Y

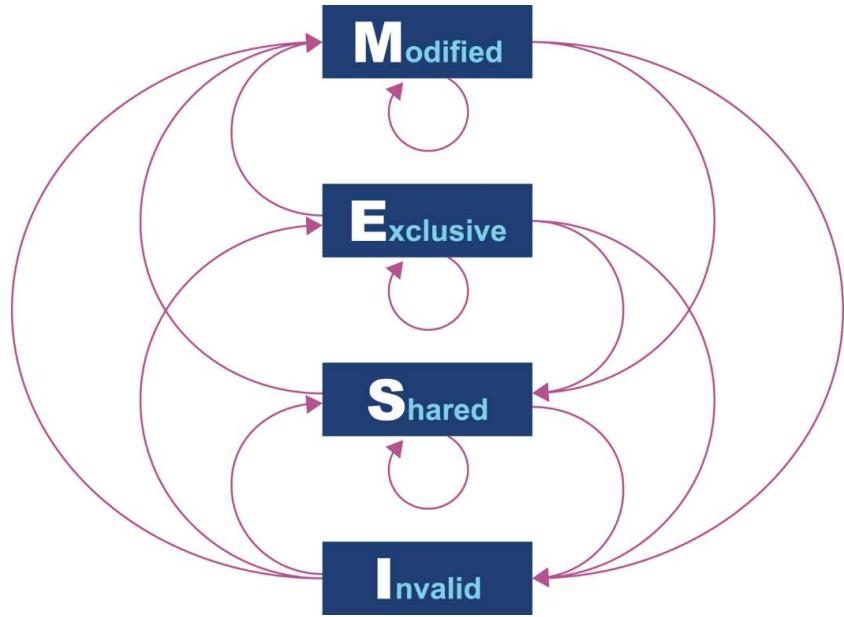
E - - - Y

S - - Y Y

I Y Y Y Y

52

Cache Consistency Protocols



53

Translation Lookaside Buffer

- > One very important use is in a different sort of cache, the Translation Lookaside Buffer (TLB).
- > TLB acts as a cache for the page tables that map virtual memory addresses to physical addresses
 - Greatly speeds up a very frequent operation
 - Access to the physical address underlying a virtual address
- > Without the TLB, all virtual address lookups would take 16 cycles, even if the page table was held in the L1 cache

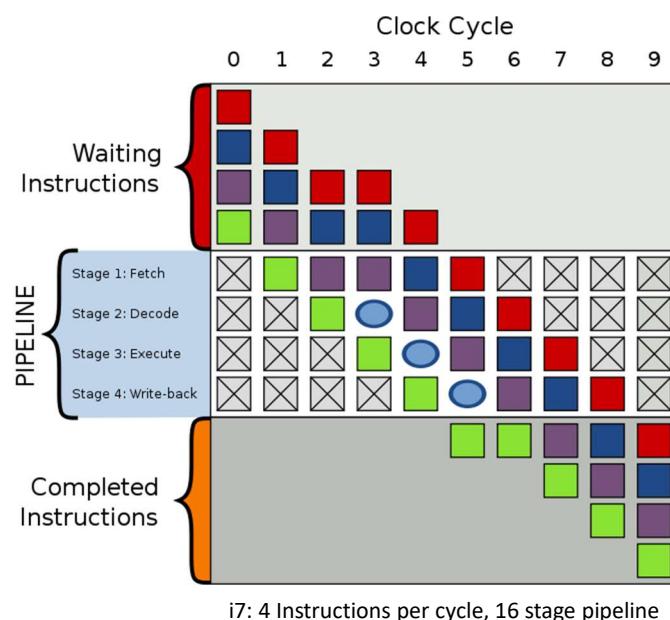
54

Branch Prediction and Speculative Execution

- > One of the advanced processor tricks that appear on modern processors is branch prediction.
- > This is used to prevent the processor from having to wait to evaluate a value needed for a conditional branch
- > Modern processors have multistage instruction pipelines
- > This means that the execution of a single CPU cycle is broken down into a number of separate stages.
- > There can be several instructions in flight (at different stages of execution) at once.

55

Branch Prediction and Speculative Execution



56

Hardware Memory Models

- > The core question about memory that must be answered in a multicore system is
“How can multiple different CPUs access the same memory location consistently?”

57

Hardware Memory Models

- > javac, JIT compiler, and CPU are all allowed to make changes to the order in which code executes.
- > This is subject to the provision that any changes don't affect the outcome as observed by the current thread
- > For example, let's suppose we have a piece of code like this:

```
myInt = otherInt;  
intChanged = true;
```

- > There is no code between the two assignments, so the executing thread doesn't need to care about what order they happen in, and thus the environment is at liberty to change the order of instructions?

58

Hardware Memory Models

	ARMv7	POWER	SPARC	x86	AMD64	zSeries
Loads moved after loads	Y	Y	-	-	-	-
Loads moved after stores	Y	Y	-	-	-	-
Stores moved after stores	Y	Y	-	-	-	-
Stores moved after loads	Y	Y	Y	Y	Y	Y
Atomic moved with loads	Y	Y	-	-	-	-
Atomic moved with stores	Y	Y	-	-	-	-
Incoherent instructions	Y	Y	Y	Y	-	Y

59

Operating Systems

- > The point of an operating system is to control access to resources that must be shared between multiple executing processes.
- > All resources are finite, and all processes are greedy
- > The need for a central system to arbitrate and meter access is essential.
- > Among these scarce resources, the two most important are usually memory and CPU time.
- > Virtual addressing via the memory management unit (MMU) and its page tables is the key feature that enables access control of memory, and prevents one process from damaging the memory areas owned by another.

60

The Scheduler

- > Access to the CPU is controlled by the process scheduler.
- > It uses a queue known as the *run queue* as a waiting area for threads or processes that are eligible to run but which must wait their turn for the CPU.
- > On a modern system there are effectively always more threads/processes that want to run than can, and so this CPU contention requires a mechanism to resolve it.
- > The job of the scheduler is to respond to interrupts, and to manage access to the CPU cores.

61

Context Switches

- > A *context switch* is the process by which the OS scheduler removes a currently running thread or task and replaces it with one that is waiting.
- > There are several different types of context switch, but broadly speaking, they all involve swapping the executing instructions and the stack state of the thread.

62

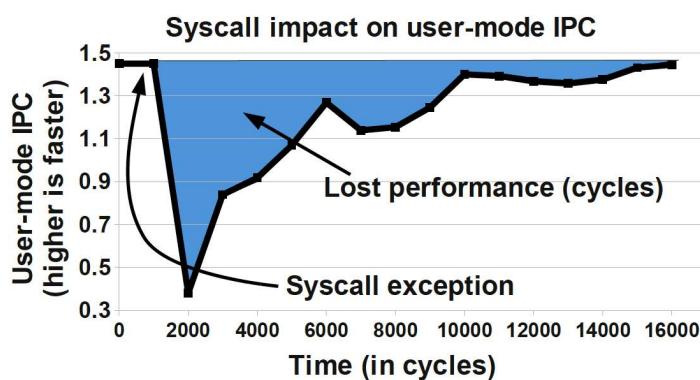
Context Switches

- > A context switch can be a costly operation, whether between user threads or from user mode into kernel mode
- > The latter case is particularly important, because a user thread may need to swap into kernel mode in order to perform some function partway through its time slice.
- > However, this switch will force instruction and other caches to be emptied, as the memory areas accessed by the user space code will not normally have anything in common with the kernel.

63

Context Switches

- > A context switch into kernel mode will invalidate the TLBs and potentially other caches.
- > When the call returns, these caches will have to be refilled, and so the effect of a kernel mode switch persists even after control has returned to user space.



64

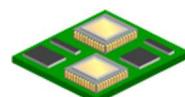
A Simple System Model

- > A simple model for describing basic sources of possible performance problems
- > The model is expressed in terms of operating system observables of fundamental subsystems and can be directly related back to the outputs of standard Unix command-line tools
- > The model is based on a simple conception of a Java application running on a Unix or Unix-like operating system.

65

Monitoring CPU Usage Overview

- > Rationale for monitoring CPU usage
 - Get big picture view of CPU demand
 - Get per process measurement of CPU utilization
- > Measurements of CPU usage
 - User (usr) time
 - System (sys) time
 - Idle time
 - Voluntary context switching (VCX)
 - Involuntary context switching (ICX)



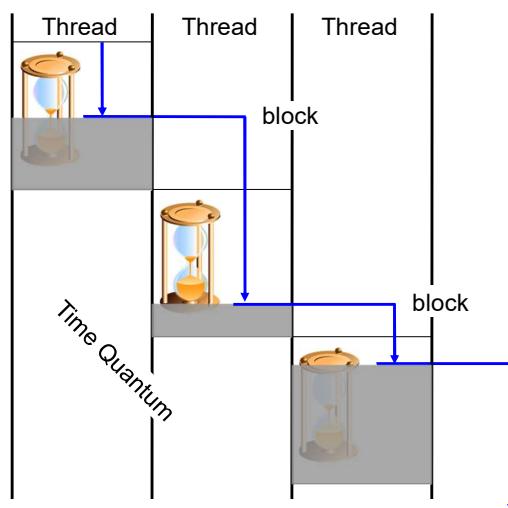
66

CPU Monitoring Performance Indicators

- > The kind of CPU statistics that may warrant further review:
 - High sys/kernel CPU time
 - Idle CPU time

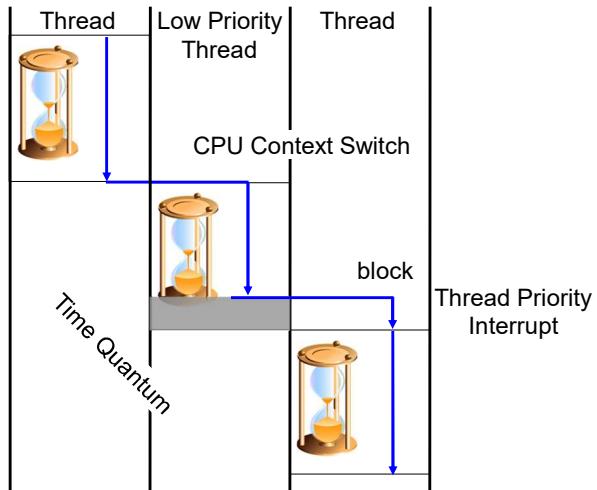
67

Voluntary Context Switching (VCX)



68

Involuntary Context Switching (ICX)



69

Tools For Monitoring CPU Usage

- > Tools to monitor CPU utilization
 - vmstat
 - mpstat
 - prstat
 - top
 - Task Manager (Windows)
 - Performance Monitor (Windows)
 - Windows Resource Manager (Windows Server)
 - Gnome System Monitor (Linux)
 - DTrace (Linux)

70

Utilizing the CPU

- > A key metric for application performance is CPU utilization.
- > CPU cycles are quite often the most critical resource needed by an application, and so efficient use of them is essential for good performance.
- > Applications should be aiming for as close to 100% usage as possible during periods of high load
- > Two basic tools that every performance engineer should be aware of are **vmstat** and **iostat**.

71

CPU Usage: vmstat

\$ vmstat 5

		memory										page			disk			faults			cpu			
kthr		r	b	w	swap	free	re	mf	pi	po	fr	de	sr	f0	s0	s1	s2	in	sy	cs	us	sy	id	
0	0	0	659620	168400	3	19	6	2	3	0	4	-0	1	-0	1	385	886	432	3	2	95			
0	0	0	129404	26456	2	87	10	0	0	0	0	5	0	0	2320	247894	4745	37	20	44				
0	0	0	127920	25524	23	395	90	0	0	0	0	27	0	1	1844	241699	4280	45	21	35				
0	0	0	126380	24204	1	92	0	0	0	0	0	0	0	0	0	2455	283529	4916	34	19	48			
0	0	0	126380	24208	1	93	0	0	0	0	0	0	0	0	0	2469	288083	4933	33	19	49			
0	0	0	126376	24204	1	79	0	0	0	0	0	0	0	0	0	2171	241732	4294	33	20	47			
0	0	0	126344	24172	1	78	1	0	0	0	0	0	16	0	3	2202	242373	4782	36	20	44			
0	0	0	126544	24372	1	85	0	0	0	0	0	0	0	0	0	2127	263456	4928	42	21	37			
0	0	0	126504	24332	1	92	0	0	0	0	0	0	0	0	0	6	2498	284583	5041	33	19	48		
0	0	0	126112	23940	1	144	2	0	0	0	0	0	0	0	0	2	2027	273269	5311	46	20	34		
0	0	0	125252	22952	3	97	53	0	0	0	0	10	0	14	2091	231099	4423	35	20	44				

72

vmstat

1. The first two columns show the number of runnable (r) and blocked (b) processes.
2. In the memory section, the amount of swapped and free memory is shown, followed by the memory used as buffer and as cache.
3. The swap section shows the memory swapped in from and out to disk (**si** and **so**).
4. The block in and block out counts (**bi** and **bo**) show the number of 512-byte blocks that have been received from and sent to a block (I/O) device.
5. The number of interrupts (**in**) and the number of context switches per second (**cs**) are displayed.

73

vmstat

6. The CPU section contains a number of directly relevant metrics, expressed as percentages of CPU time. In order, they are user time (us), kernel time (sy, for “system time”), idle time (id), waiting time (wa), and “stolen time” (st, for virtual machines).

74

CPU Usage

- > if we observe that the CPU utilization is not approaching 100% user time, then the next obvious question is, "Why not?"
- > What is causing the program to fail to achieve that?
- > Are involuntary context switches caused by locks the problem?
- > Is it due to blocking caused by I/O contention?

75

CPU Usage: mpstat

> \$ **mpstat 5**

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	idle
4	0	0	0	101	100	10	0	0	0	0	1	0	1	0	99		
5	131	0	550	105	100	1246	4	0	1	0	4877	42	10	0	48		
8	1	0	0	101	100	53	0	1	0	0	355	0	1	0	98		
9	23	0	0	102	100	363	0	0	0	0	1283	3	3	0	94		
10	25	0	0	101	100	442	1	0	0	0	1426	3	4	0	93		
11	0	0	0	2135	2134	34	0	0	0	0	15	0	8	0	92		
12	0	0	0	101	100	35	0	1	1	0	4	0	1	0	99		
13	0	0	0	101	100	4	0	0	0	0	1	0	1	0	99		
14	0	0	62	102	101	8	0	0	0	0	1	0	1	0	99		
15	0	0	0	101	100	9	0	0	0	0	0	0	2	0	98		
CPU	minf	mjf	xcal	intr	ithr	csw icsw	migr	smtx	srw	syscl	usr	sys	wt	idl			
0	0	0	1107	400	300	15	0	0	0	0	7	0	2	0	98		
1	0	0	0	101	100	13	0	0	0	0	0	0	1	0	99		
4	0	0	0	101	100	11	0	0	0	0	1	0	1	0	99		
5	46	0	470	103	100	117	2	0	0	0	1240	22	4	0	75		
8	0	0	0	107	100	35	6	1	0	0	65	73	1	0	26		
9	1	0	0	102	101	31	0	1	0	0	114	0	1	0	99		
10	0	0	0	101	100	17	0	1	0	0	125	0	1	0	99		
11	0	0	0	217	216	12	0	0	0	0	26	0	1	0	99		
12	0	0	0	102	101	58	0	0	1	0	45	0	1	0	99		
13	0	0	0	102	101	5	0	0	0	0	0	0	1	0	99		
14	0	0	0	102	101	6	0	0	0	0	1	0	1	0	99		
15	0	0	0	101	100	9	0	0	0	0	1	0	1	0	99		

76

CPU Usage: prstat

\$ prstat

PID	USERNAME	SIZE	RSS	STATE	PRI	NICE	TIME	CPU	PROCESS/NLWP
2647	huntrch	674M	632M	cpu1	50	0	0:08:02	38%	java/18
778	huntrch	49M	23M	sleep	59	0	0:00:25	0.3%	Xorg/1
1136	root	103M	22M	sleep	59	0	0:00:35	0.1%	java/19
2264	huntrch	3924M	1928K	sleep	59	0	0:00:01	0.1%	cpubar/1
1122	noaccess	76M	15M	sleep	59	0	0:00:32	0.1%	java/24
1708	huntrch	110M	9800K	sleep	59	0	0:00:24	0.1%	mixer_applet2/1
525	root	3948K	392K	sleep	59	0	0:00:09	0.1%	vmware-guestd/1
2251	huntrch	55M	10M	sleep	59	0	0:00:15	0.0%	gnome-terminal/2
2653	huntrch	4636K	1096K	cpu0	59	0	0:00:00	0.0%	prstat/1
1198	huntrch	51M	752K	sleep	59	0	0:00:20	0.0%	vmware-user/1
2841	huntrch	47M	10M	sleep	59	0	0:00:00	0.0%	screenshot/1
749	root	4544K	1864K	sleep	59	0	0:00:01	0.0%	nsqd/23
1702	huntrch	107M	9596K	sleep	59	0	0:00:02	0.0%	clock-applet/1
9	root	8728K	1252K	sleep	59	0	0:00:38	0.0%	svc.configd/13
1589	huntrch	110M	11M	sleep	59	0	0:00:01	0.0%	gnome-panel/1
604	root	6232K	2568K	sleep	59	0	0:00:01	0.0%	intrd/1
2840	huntrch	42M	4696K	sleep	59	0	0:00:00	0.0%	script-fu/1
2656	huntrch	64M	24M	sleep	59	0	0:00:06	0.0%	gimp-2.3/5
2247	root	4760K	1364K	sleep	59	0	0:00:01	0.0%	devfsadm/8
1592	huntrch	140M	21M	sleep	59	0	0:00:09	0.0%	nautilus/1
1646	huntrch	14M	3888K	sleep	59	0	0:00:00	0.0%	gnome-vfs-daemo/2
Total: 80 processes, 254 lwps, load averages: 5.45, 3.08, 1.33									

77

CPU Usage: prstat -m

PID	USERNAME	USR	SYS	TRP	TFL	DFL	LCK	SLP	LAT	VCX	TCX	SCL	SIG	PROCESS/LWPID	
6506	huntrch	49	24	0.2	0.0	0.0	1.1	19	7.3	4K	1K	.13	0	java/2	
6505	huntrch	11	7.1	0.0	0.0	0.0	0.3	74	7.5	4K	162	87K	0	java/2	
766	huntrch	1.2	0.5	0.0	0.0	0.0	0.0	94	2.1	266	131	1K	64	Xorg/1	
6506	huntrch	1.2	0.6	0.0	0.0	0.0	0.0	98	0.0	0.3	595	10	1K	0	java/3
990	huntrch	1.1	0.1	0.0	0.0	0.0	0.0	97	1.5	8	64	291	0	wnck-applet/1	
6506	huntrch	0.5	0.1	0.0	0.0	0.0	0.0	97	0.0	2.3	195	3	196	0	java/5
6506	huntrch	0.3	0.2	0.0	0.0	0.0	0.0	99	0.0	0.1	193	7	385	0	java/4
972	huntrch	0.3	0.2	0.0	0.0	0.0	0.0	99	0.2	97	0	549	0	metacity/1	
6190	huntrch	0.2	0.3	0.0	0.0	0.0	0.0	99	0.5	165	1	647	0	java/21	
4919	huntrch	0.4	0.0	0.0	0.0	0.0	0.0	100	0.1	27	0	78	0	soffice.bin/1	
6417	huntrch	0.3	0.0	0.0	0.0	0.0	0.0	100	0.0	4	0	8	0	thunderbird/-1	
6505	huntrch	0.2	0.1	0.0	0.0	0.0	0.0	100	0.0	0.0	44	0	204	0	java/3
6512	huntrch	0.2	0.1	0.0	0.0	0.0	0.0	100	0.0	17	23	111	0	screenshot/1	
6506	huntrch	0.1	0.1	0.0	0.0	0.0	0.0	99	0.4	151	38	87	0	java/9	
1095	huntrch	0.1	0.1	0.0	0.0	0.0	0.0	100	0.0	0.5	33	0	211	0	firefox-bin/1
6190	huntrch	0.1	0.1	0.0	0.0	0.0	0.0	99	0.4	156	24	87	0	java/14	
6505	huntrch	0.2	0.0	0.0	0.0	0.0	0.0	99	0.0	0.6	32	0	32	0	java/5
980	huntrch	0.2	0.0	0.0	0.0	0.0	0.0	99	0.7	12	1	46	0	nautilus/1	
6494	huntrch	0.2	0.0	0.0	0.0	0.0	0.0	100	0.0	13	0	58	0	gimp-2.4/1	
6505	huntrch	0.1	0.1	0.0	0.0	0.0	0.0	100	0.2	156	1	87	0	java/9	
1095	huntrch	0.1	0.1	0.0	0.0	0.0	0.0	100	0.0	0.2	104	0	150	0	firefox-bin/3
Total: 99 processes, 321 lwps, load averages: 1.60, 1.28, 0.64															

78

Monitoring Network I/O Overview

- > Data of interest
 - Network utilization in terms of Transaction Control Protocol (TCP) statistics and established connections
- > Tools to monitor network I/O
 - **netstat** (Linux)
 - Performance Monitor (Windows)
 - **DTrace** (Linux)
 - **tcptop** (DTrace Toolkit)



79

Network I/O: Using **tcptop**

- > **tcptop** can show per process TCP statistics:

```
huntrch@ditka: ~
File Edit View Terminal Tabs Help
# tcptop -C 10
Sampling... Please wait.
2005 Jul 5 04:55:25, load: 1.11, TCPin: 2 Kb, TCPout: 110 Kb

UID PID LADDR      LPORT FADDR      FPORT      SIZE NAME
100 20876 192.168.1.5 36396 192.168.1.1    79      1160 finger
100 20875 192.168.1.5 36395 192.168.1.1    79      1160 finger
100 20878 192.168.1.5 36397 192.168.1.1    23      1303 telnet
100 20877 192.168.1.5 859 192.168.1.1    514     115712 rcp
```

The screen capture shows **rcp** generating 115 kb of traffic.

80

Network I/O: Using **nicstat**

> **nicstat 1**

Time	Int	rKb/s	wKb/s	rPk/s	wPk/s	rAvs	wAvs	%Util	Sat
12:33:04	hme0	1.51	4.84	7.26	10.32	213.03	480.04	0.05	0.00
12:33:05	hme0	0.20	0.26	3.00	3.00	68.67	90.00	0.00	0.00
12:33:06	hme0	0.14	0.26	2.00	3.00	73.00	90.00	0.00	0.00
12:33:07	hme0	0.14	0.52	2.00	6.00	73.00	88.00	0.01	0.00
12:33:08	hme0	0.24	0.36	3.00	4.00	81.33	92.00	0.00	0.00
12:33:09	hme0	2.20	1.77	16.00	18.00	140.62	100.72	0.03	0.00
12:33:10	hme0	0.49	0.58	8.00	9.00	63.25	66.00	0.01	0.00
12:33:11	hme0	12.16	1830.38	185.06	1326.42	67.26	1413.06	15.09	0.00
12:33:12	hme0	19.03	3094.19	292.88	2229.11	66.53	1421.40	25.50	0.00
12:33:13	hme0	19.55	3151.87	301.00	2270.98	66.50	1421.20	25.98	0.00
12:33:14	hme0	11.99	1471.67	161.07	1081.45	76.25	1393.49	12.15	0.00
12:33:15	hme0	0.14	0.26	2.00	3.00	73.00	90.00	0.00	0.00
12:33:16	hme0	0.14	0.26	2.00	3.00	73.00	90.00	0.00	0.00
12:33:17	hme0	0.14	0.26	2.00	3.00	73.00	90.00	0.00	0.00

81

Monitoring Disk I/O Overview

- > Data of interest
 - Number of disk accesses
 - Latency and average latencies
- > Tools to monitor disk I/O
 - **iostat** (Linux)
 - **iotop** (Linux)
 - **pidstat** (Linux)
 - Performance Monitor (Windows)
 - **DTrace**
- > Disk caches



82

Disk I/O: iotop

- > **iotop** reporting at a 5-second interval
- > **DISKTIME** reported in microseconds

A screenshot of a terminal window titled "huntrch@ditka: ~". The window displays the output of the "iotop" command. The output shows system load (load: 1.10), disk reads (disk_r: 5302 Kb), and disk writes (disk_w: 20 Kb). It lists processes by UID, PID, PPID, CMD, DEVICE, MAJ, MIN, D, and DISKTIME. A red box highlights the "DISKTIME" column for the "find" process, which has a value of 3094794 microseconds.

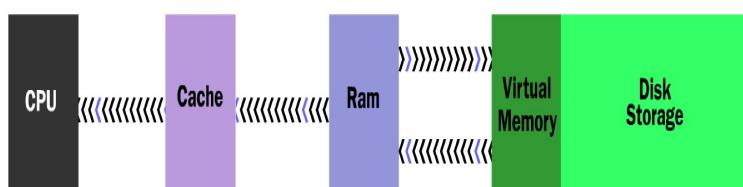
UID	PID	PPID	CMD	DEVICE	MAJ	MIN	D	DISKTIME
0	0	0	sched	cmdk0	102	0	W	532
0	0	0	sched	cmdk0	102	0	R	245398
0	27758	20320	find	cmdk0	102	0	R	3094794

83

Monitoring Virtual Memory: Overview

- > Observe paging to identify swapping
 - Pages in (pi)
 - Pages out (po)
 - Scan rate (sr)
- > Fixing the swapping problem
- > Why is swapping bad for a Java application?
 - Want to explain? (Without looking at your notes)

Memory Management



84

Virtual Memory Tools

- > Tools to monitor memory paging and usage
 - **vmstat** (Linux)
 - **top**
 - Performance Monitor (Windows)
 - **DTrace** (Linux)

85

Virtual Memory: vmstat

\$ **vmstat 5**

kthr	memory	page	disk	faults	cpu
r b w	swap free re mf pi po fr de sr m0 m1 m2 m3	in sy cs us sy id			
0 0 0	11532280 3322672 4 18 6 0 0 0 0 0 1 1 1 0 1599 491	147 1 1 98			
0 0 0	11406040 3186248 0 7 0 0 0 0 0 0 0 0 0 0 1567 348	138 0 1 99			
0 0 0	11406832 3187048 0 1 0 0 0 0 0 19 19 19 0 1737 332	153 0 1 99			
0 0 0	11408128 3188344 0 1 0 2 2 0 0 0 0 0 0 0 1573 336	134 0 1 99			
0 0 0	11407912 3188128 0 1 0 0 0 0 0 0 0 0 0 0 1567 331	139 0 1 99			
0 0 0	11407464 3187696 2 21 0 0 0 0 0 17 17 17 0 1717 331	157 0 1 99			
0 0 0	11407472 3187704 0 1 0 0 0 0 0 0 0 0 0 0 1563 310	149 0 1 99			
0 0 0	11407520 3187744 0 2 0 0 0 0 0 0 0 0 0 0 1565 317	144 0 1 99			
0 0 0	11407504 3187720 0 1 0 0 0 0 0 0 0 0 0 0 1561 305	138 0 1 99			
0 0 0	11407016 3187216 32 196 0 0 0 0 0 0 0 0 0 0 2764 4997 1381	1 2 97			
0 0 0	11396280 3176232 43 351 0 0 0 0 0 0 0 0 0 0 5036 14128 3772	9 4 87			
0 0 0	11399760 3179872 25 65 0 0 0 0 0 8 8 8 0 1740 2450	314 16 2 82			
0 0 0	11407392 3187600 0 1 0 0 0 0 0 0 0 0 0 0 1579 332	150 17 1 82			
0 0 0	11407352 3187568 0 3 0 0 0 0 0 0 0 0 0 0 1572 313	138 17 1 82			
0 0 0	11407312 3187536 0 1 0 0 0 0 0 0 0 0 0 0 1573 324	143 12 1 86			
0 0 0	11407280 3187504 0 1 0 2 2 0 0 0 0 0 0 0 1564 316	142 0 1 99			
0 0 0	11407248 3187464 0 1 0 0 0 0 0 0 0 0 0 0 1562 319	147 0 1 99			
0 0 0	11407216 3187440 0 1 0 0 0 0 0 3 3 3 0 1593 306	140 0 1 99			
kthr	memory	page	disk	faults	cpu
r b w	swap free re mf pi po fr de sr m0 m1 m2 m3	in sy cs us sy id			
0 0 0	11407408 3187712 0 4 0 0 0 0 0 0 0 0 0 1583 557	179 0 1 98			

86

Virtual Memory: Swapping Example

hutch@ditka: ~																		
File Edit View Terminal Tabs Help																		
kthr	memory			page			disk			faults			cpu					
r b w	swap	free	re	mf	pi	po	fr	de	srf	f0	s0	s1	s2	in	sy	cs	us	sy id
1 0 0	499792	154720	1	1697	0	0	0	0	0	0	0	0	0	12	811	612	1761	90 7 4
1 0 0	498856	44052	1	3214	0	0	0	0	0	0	0	0	0	12	1290	2185	3078	66 18 15
3 0 0	501188	17212	1	1400	2	2092	4911	0	37694	0	53	0	12	5262	3387	1485	52 27 21	
1 0 0	500696	20344	26	2562	13	4265	7553	0	9220	0	66	0	12	1192	3007	2733	71 17 12	
1 0 0	499976	20108	3	3146	24	3032	10009	0	10971	0	63	0	6	1346	1317	3358	78 15 7	
1 0 0	743664	259080	61	1706	70	8882	10017	0	19866	0	178	0	52	1213	595	688	70 12 18	

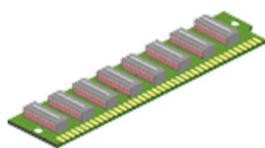
> Data of interest

- pi – pages in; po – pages out; sr – page scan rate
- Watch for high scan rate (see rows 3 to 6), or increasing trend. Low scan rate is ok if they occur infrequently.

87

Virtual Memory: Fixing Swapping Problem

- > Add physical memory.
- > Reduce number of applications running on the machine.
- > Anyone, or any combination of the above will help.



88

Monitoring Processes Overview

- > Data of interest
 - Footprint size
 - Number of threads and thread state
 - CPU usage
 - Runtime stack
 - Context switches
 - Lock contention

89

Process Monitoring Tools

- > **ps** (Linux)
- > **vmstat** (Linux)
- > **mpstat** (Linux)
- > **pidstat** (Linux)
- > Performance Monitor (Windows)
- > **top**
- > **DTrace** (Linux)

90

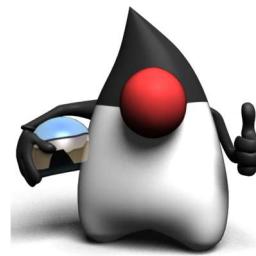
Processes: mpstat

huntrch@ditka: ~																		
CPU	mif	mif	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl			
0	21	1	6	319	133	261	18	23	3	0	551	4	3	0	92			
1	19	1	4	36	14	59	15	23	3	0	491	4	3	0	93			
2	17	1	4	19	12	61	13	23	3	0	355	4	3	0	90			
3	18	1	4	16	15	49	12	23	3	0	390	4	3	0	91			
CPU	mif	mif	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl			
0	28	2	0	192	83	92	32	14	2	0	185	78	15	0	7			
1	49	1	0	37	1	80	28	16	2	0	139	80	16	0	4			
2	28	1	0	20	7	94	34	17	1	0	283	83	12	0	5			
3	39	1	2	52	1	99	36	16	3	0	219	74	19	0	7			
CPU	mif	mif	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl			
0	34	0	2	171	75	78	32	12	1	0	173	90	9	0	2			
1	38	1	0	39	1	84	29	13	2	0	153	66	12	0	23			
2	28	8	0	21	9	97	31	20	2	0	167	67	13	0	20			
3	35	3	1	43	1	98	29	20	3	0	190	52	25	0	23			

91

Monitoring the Kernel

- > Data of interest
 - Kernel CPU utilization, locks, system calls, interrupts, migrations, run queue depth
- > Tools to monitor the kernel
 - vmstat (Linux)
 - mpstat (Linux)
 - Performance Monitor (Windows)
 - DTrace (Linux)



92

Kernel: vmstat

procs			memory						swap		io		system				cpu			
r	b	w	swpd	free	buff	cache	si	so	bi	bo	in	cs	us	sy	id	wa				
4	0	0	0	959476	340784	1387176	0	0	0	0	1030	8977	63	35	1	0				
3	0	0	0	959444	340784	1387176	0	0	0	0	1014	7981	62	36	2	0				
6	0	0	0	959460	340784	1387176	0	0	0	16	1019	9380	63	36	1	0				
1	0	0	0	958820	340784	1387176	0	0	0	0	1036	9157	63	35	2	0				
4	0	0	0	958500	340784	1387176	0	0	0	29	1012	8582	62	37	1	0				

- > The “us” column shows the percentage of user CPU utilization.
- > The “sy” column shows the percentage of kernel or system CPU utilization.
- > The “id” column shows the percentage of idle or available CPU.
- > The sum of the “us” column and “sy” column should be equal to 100 minus the value in the “id” column, that is, $100 - (\text{id} \text{ column value})$.

93

Kernel: vmstat

kthr			memory						page						disk						faults				cpu			
r	b	w	swpd	free	re	mf	pi	po	fr	de	sr	f0	s0	s1	s2	in	sy	cs	us	sy	id							
0	0	0	672604	141500	10	40	36	6	10	0	20	0	3	0	2	425	1043	491	4	3	93							
1	1	0	888460	632992	7	32	97	0	0	0	0	0	21	0	12	462	1099	429	32	19	49							
0	1	0	887848	631772	4	35	128	0	0	0	0	0	30	0	13	325	575	314	38	13	49							
0	1	0	887592	630844	6	26	79	0	0	0	0	0	40	0	11	324	501	287	36	10	54							
1	0	0	887304	630160	5	33	112	0	0	0	0	0	50	0	16	369	899	367	37	11	52							
0	1	0	886920	629092	4	30	101	0	0	0	0	0	26	0	18	354	707	260	39	14	46							

94

Kernel: mpstat

hunth@ditka: ~																	
CPU		minf	mjf	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl	
0	21	1	6	319	133	261	18	23	3	0	551	4	3	0	92		
1	19	1	4	36	14	59	15	23	3	0	491	4	3	0	93		
2	17	1	4	19	12	61	13	23	3	0	355	4	3	0	90		
3	18	1	4	16	15	49	12	23	3	0	390	4	3	0	91		
CPU		minf	mjf	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl	
0	28	2	0	192	83	92	32	14	2	0	185	78	15	0	7		
1	49	1	0	37	1	80	28	16	2	0	139	80	16	0	4		
2	28	1	0	20	7	94	34	17	1	0	283	83	12	0	5		
3	39	1	2	52	1	99	36	16	3	0	219	74	19	0	7		
CPU		minf	mjf	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl	
0	34	0	2	171	75	78	32	12	1	0	173	90	9	0	2		
1	38	1	0	39	1	84	29	13	2	0	153	66	12	0	23		
2	28	8	0	21	9	97	31	20	2	0	167	67	13	0	20		
3	35	3	1	43	1	98	29	20	3	0	190	52	25	0	23		

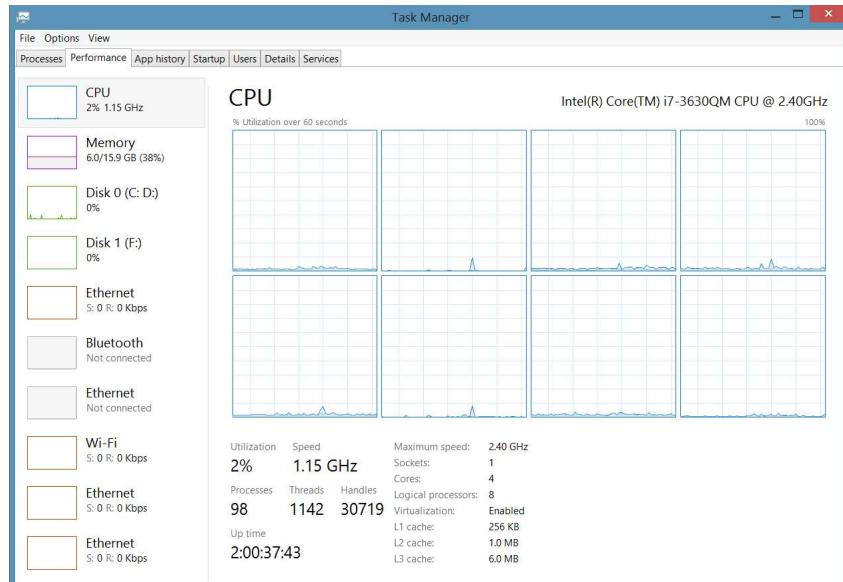
95

Kernel: mpstat

CPU	minf	mjf	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl
0	28	2	0	192	83	92	32	14	2	0	185	78	15	0	7
1	49	1	0	37	1	80	28	16	2	0	139	80	16	0	4
2	28	1	0	20	7	94	34	17	1	0	283	83	12	0	5
3	39	1	2	52	1	99	36	16	3	0	219	74	19	0	7
CPU	minf	mjf	xcal	intr	ithr	csw	icsw	migr	smtx	srw	syscl	usr	sys	wt	idl
0	34	0	2	171	75	78	32	12	1	0	173	90	9	0	2
1	38	1	0	39	1	84	29	13	2	0	153	66	12	0	23
2	28	8	0	21	9	97	31	20	2	0	167	67	13	0	20
3	35	3	1	43	1	98	29	20	3	0	190	52	25	0	23

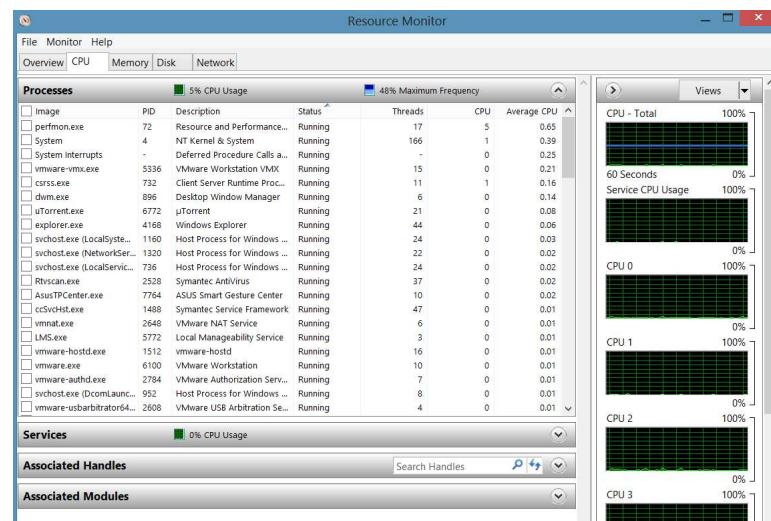
96

Monitoring CPU Utilization on Windows

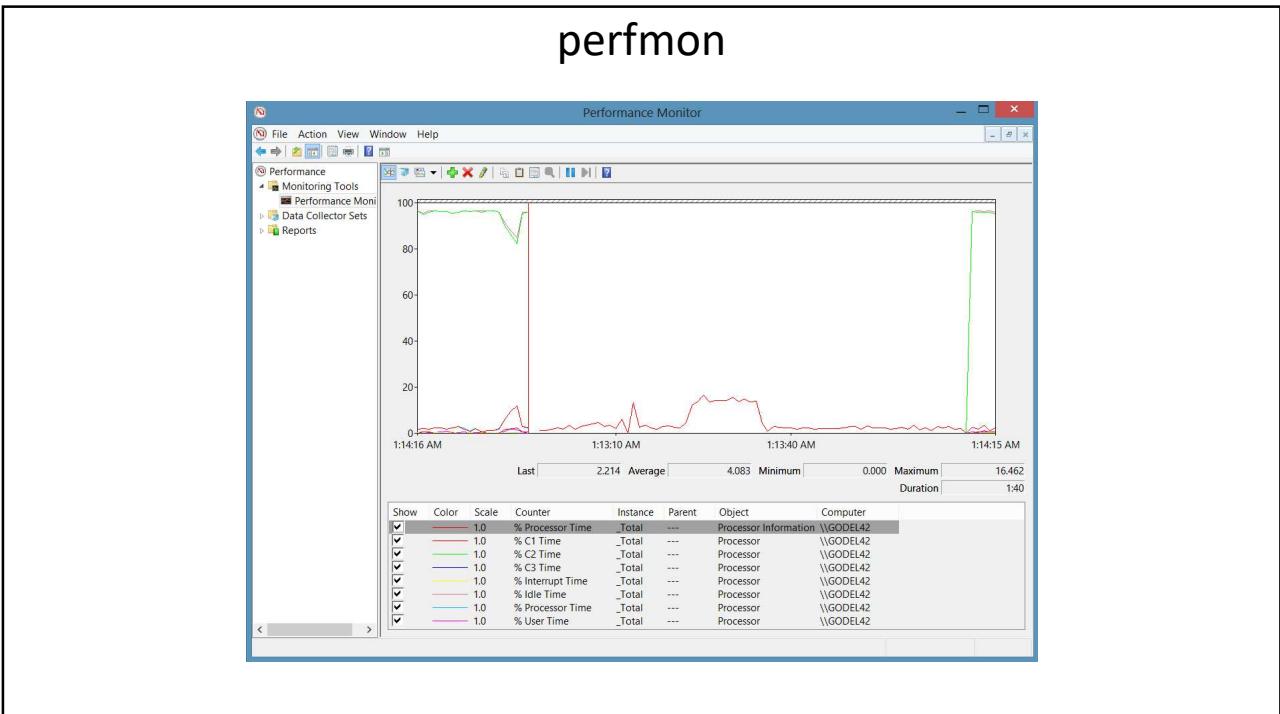


97

Monitoring CPU Utilization on Windows



98



99

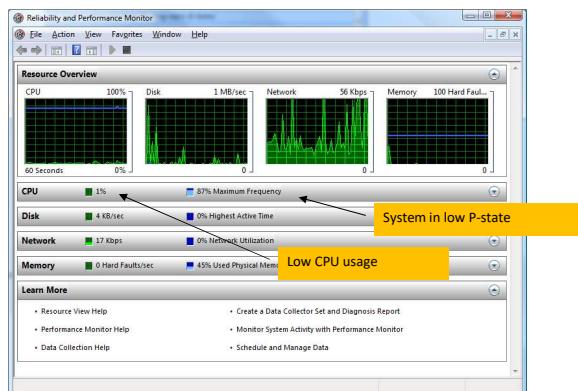
Processor Power Management in Windows

- > Windows OS includes support for ACPI processor power management (PPM) features
 - support for processor performance states
 - support for processor idle sleep states on multiprocessor systems.
 - > PPM technologies are defined in the ACPI specification and are divided into three categories:
 - Processor performance states (ACPI “P-states”)
 - Linear stop clock throttle states (ACPI “T-states”)
 - Processor idle sleep states (ACPI “C-states”)
 - > For Intel, driver name is **Intelppm.sys**
- powercfg -qh scheme_current sub_processor**

100

resmon

- > Windows Reliability and Performance Monitor Showing Performance State Usage



101

typeperf

- > typeperf "\Processor(_Total)\% Privileged Time" "\Processor(_Total)\% User Time"

```
C:\Windows\system32\typeperf.exe

"(PDH-CSU 4.0)", "\GODEL42\Processor(_Total)\% Privileged Time", "\GODEL42\Processor(_Total)\% User Time"
"03/23/2013 01:16:56.461", "0.000000", "0.584963"
"03/23/2013 01:16:57.463", "0.584956", "0.194985"
"03/23/2013 01:16:58.464", "0.389980", "1.364914"
"03/23/2013 01:16:59.466", "0.194985", "0.194985"
"03/23/2013 01:17:00.468", "0.194985", "0.389970"
"03/23/2013 01:17:01.469", "0.194995", "0.000000"
"03/23/2013 01:17:02.471", "0.584955", "0.000000"
"03/23/2013 01:17:03.473", "0.974943", "0.000000"
"03/23/2013 01:17:04.474", "0.974934", "0.584965"
"03/23/2013 01:17:05.476", "0.194985", "0.389980"
"03/23/2013 01:17:06.478", "0.194985", "0.000000"
"03/23/2013 01:17:07.479", "0.779950", "0.779950"
"03/23/2013 01:17:08.481", "0.974944", "0.194985"
"03/23/2013 01:17:09.483", "0.389970", "0.389970"
"03/23/2013 01:17:10.484", "0.389980", "0.000000"
"03/23/2013 01:17:11.486", "1.169893", "0.000000"
"03/23/2013 01:17:12.488", "0.000000", "0.000000"
"03/23/2013 01:17:13.489", "0.779950", "0.000000"
"03/23/2013 01:17:14.491", "0.000000", "0.194995"
"03/23/2013 01:17:15.493", "0.194985", "0.000000"
"03/23/2013 01:17:16.494", "0.194995", "0.194985"
```

102

typeperf

- > You can also assemble a list of performance counters in a file and pass the name of the file to the typeperf command.
- > For example, you can enter the following performance counters in a file named cpu-util.txt:

```
\Processor(_Total)\% Privileged Time  
\Processor(_Total)\% User Time
```

- > Then, invoke the **typeperf** command with the option **-cf** followed by the file name.

```
typeperf -cf cpu-util.txt
```

103

logman

- > Helps to manage *Performance Monitor* and *performance* logs from the command line.

```
cmd> set _mycnt="\Processor(_Total)\% Processor Time"  
cmd> set _mylogfile=C:\tmp\ss64.blg  
cmd> logman create counter ss64-CPU -f bincirc -v mmddhhmm  
      -max 250 -c %_mycnt% -o %_mylogfile%
```

The command completed successfully.

```
cmd> logman start ss64-CPU
```

The command completed successfully.

```
cmd> logman stop ss64-CPU
```

The command completed successfully.

104



OPTIMIZATION

105

Memory Management

- > Preallocate Memory
 - Use `array`, `np.array`, or custom buffers for predictable memory usage and faster access.
- > Minimize Object Creation
 - Reduce unnecessary object creation with loops and comprehensions.
 - Favor `itertools` functions.
- > Properly Handle Closures
 - Closures capture references, potentially causing memory leaks.
 - Use `nonlocal` or `functools.partial` when needed.
- > Profile for Memory Bottlenecks
 - Use tools like `memory_profiler` to identify areas with high object churn and target them for optimization.

106

Memory Management

- > Garbage collection tuning
 - Tune garbage collector settings based on your application's memory usage patterns.
 - Consider tools like Pympler to analyze memory allocation.
- > Close file handles and connections
 - Ensure proper resource management by closing files, database connections, and network sockets explicitly to avoid memory leaks.
- > Minimize object creation
 - Avoid creating unnecessary temporary objects within loops.
 - Consider object pooling for frequently used objects.

107

Memory Management

- > Use weak references
 - Use `weakref` to hold non-critical references to objects, allowing them to be garbage collected without impacting others.

108

Data Structures and Algorithms

- > Choose the Right Data Structure
 - Use sets for membership checks, dictionaries for fast key-value access, and lists for ordered sequences.
- > Algorithm Efficiency Matters
 - Utilize appropriate sorting algorithms like Timsort or quicksort for large datasets.
- > Avoid Nested Loops
 - Unroll inner loops where possible for clearer code and potential performance gains.
- > Memoization for Repeated Calculations
 - Cache expensive function calls with techniques like `functools.cache` for repeated inputs.

109

Code Structure and Optimization

- > Utilize Comprehensions
 - Use list/dictionary comprehensions instead of explicit loops for concise and often faster code.
- > Vectorize Operations
 - Leverage NumPy for vectorized operations on large arrays, significantly boosting performance.
- > Optimize Imports
 - Organize and defer imports strategically to avoid unnecessary module loading at runtime.
- > Reduce Function Calls
 - Inline short functions within loops to minimize overhead and improve efficiency.

110

External Factors and Tools

- > Profile Your Code
 - Use profilers like cProfile or hotshot to identify execution hotspots and target bottlenecks.
- > Utilize Caching Mechanisms
 - Cache database queries, API calls, or expensive computations for better responsiveness.
- > Consider Asynchronous Programming
 - Utilize non-blocking I/O libraries like asyncio for improved concurrency and responsiveness.
- > Monitor System Resources
 - Track CPU, memory, and network usage to identify potential bottlenecks and resource limitations.

111

Continuous Improvement

- > Benchmark Regularly
 - Track performance changes as your code evolves and identify regressions or potential improvements.
- > Readability Matters
 - Prioritize maintainable and clear code while optimizing, even if it leads to slightly slower execution.
- > Stay Updated
 - Utilize modern Python features and libraries for optimized performance and improved syntax.
- > Seek Community Resources
 - Leverage online resources, forums, and profiling tools to learn from others and improve your optimization skills.

112

Advanced Techniques

- > Cythonize critical sections
 - Translate performance-critical parts of your code to C for significant speedups.
- > Numba for JIT compilation
 - Use Numba to JIT compile specific functions for improved performance of numerical code.
- > Multithreading and multiprocessing
 - Leverage multiple cores or CPUs for parallel execution of independent tasks.
 - Be mindful of overhead and data synchronization challenges.
- > Asynchronous programming
 - Use libraries like asyncio for non-blocking I/O operations and improve responsiveness, especially for network-bound applications.

113

Advanced Techniques

- > Memory-mapped files
 - Share data directly between Python and C++ code for efficient data exchange and avoid unnecessary copying

114

Database Optimization

- > Optimize queries
 - Use proper indexing, avoid unnecessary SELECT * statements, and utilize JOINs efficiently.
 - Consider database-specific query optimization strategies.
- > Caching database results
 - Cache frequently used database queries to reduce server load and improve response times.
- > Database connection pooling
 - Reduce overhead by maintaining a pool of pre-established database connections.

115

Testing and Monitoring

- > Write unit tests for optimizations
 - Ensure your optimizations haven't introduced regressions.
 - Test both performance and functionality.
- > Monitor application performance
 - Use tools like New Relic or Prometheus to track resource usage, identify bottlenecks, and measure the impact of optimization efforts.
- > Continuous integration and deployment
 - Automate performance testing and optimization changes into your CI/CD pipeline for a proactive approach.

116

Introduction to Profiling

- > What is Profiling?
 - Analyzing program performance
 - Identifying bottlenecks in the code
- > Why Profile Your Code?
 - Optimize execution time
 - Improve resource usage
 - Enhance user experience

117

Benchmarking and Profiling

- > Recognizing the slow parts of your program is the single most important task when it comes to speeding up your code
- > When designing a performance-intensive program, the very first step is to write your code without bothering with small optimizations:

"Premature optimization is the root of all evil."

Donald Knuth

118

Benchmarking and Profiling

- > Make it run
 - Ensure that it produces the correct results.
- > Make it right
 - Ensure that the design of the program is solid.
- > Make it fast
 - Once our program is working and is well structured, you can focus on performance optimization.

119

What is cProfile?

- > cProfile is Python's built-in deterministic profiler.
- > Provides detailed statistics about program execution.
- > Advantages
 - Lightweight
 - Easy to use
 - Highly integrated with Python's standard library

120

Installing and Setting Up cProfile

- > Pre-installed with Python
 - No installation needed.
- > Import directly

```
import cProfile
```

 - Use via command line

```
python -m cProfile [script.py]
```

121

Key Features of cProfile

- > Execution Profiling
 - Tracks function calls
- > Detailed Statistics
 - Call counts, execution times, etc.
- > Output Formats
 - Plain text
 - Binary (for visualization tools)
- > Ease of Integration
 - Works seamlessly with Python programs

122

Benchmarking with pytest-benchmark

```
from simul import Particle, ParticleSimulator
def test_evolve(benchmark):
    # ... previous code
    benchmark(simulator.evolve, 0.1)
```

```
=====
platform linux -- Python 3.5.2, pytest-3.0.5, py-1.4.32, pluggy-0.4.0
benchmark: 3.0.0 (defaults: timer=time.perf_counter disable_gc=False min_rounds=5 min_time=5.00us max_time=1.00s calibration_precision=10 warmup=False warmup_iterations=100000)
rootdir: /home/gabriele/workspace/hyperf/chapter1, inifile:
plugins: benchmark-3.0.0
collected 2 items

test_simul.py .

----- benchmark: 1 tests -----
Name (time in ms)      Min     Max     Mean   StdDev   Median      IQR  Outliers(*)  Rounds Iterations
test_evolve      29.4716  41.1791  30.4622  2.0234  29.9630  0.7376          2;2       34           1
----- (*) Outliers: 1 Standard Deviation from Mean; 1.5 IQR (InterQuartile Range) from 1st Quartile and 3rd Quartile.
===== 1 passed in 2.52 seconds =====
```

123

Finding bottlenecks with cProfile

- > Two profiling modules are available through the Python standard library
- > The `profile` module
 - This module is written in pure Python and adds a significant overhead to the program execution.
 - Its presence in the standard library is because of its vast platform support and because it is easier to extend.
- > The `cProfile` module
 - This is the main profiling module, with an interface equivalent to `profile`.
 - It is written in C, has a small overhead, and is suitable as a general-purpose profiler.

124

Using cProfile: Command Line

- > Basic Command
 - `python -m cProfile script.py`
- > Sort Options
 - `-s time`
 - Sort by execution time
 - `-s calls`
 - Sort by call count
- > Saving Results:
`python -m cProfile -o output.prof script.py`

125

cProfile

- > cProfile does not require any change in the source code and can be executed directly on an existing Python script or function.
- > You can use cProfile from the command line in this way:
`python -m cProfile simul.py`
`python -m cProfile -s tottime simul.py`
`python -m cProfile -o prof.out simul.py`

126

Using cProfile: Programmatically

- > The usage of cProfile as a Python module requires invoking the cProfile.run function in the following way:

```
import cProfile
def fun():
    # function logic

cProfile.run("fun()")
```

127

cProfile

- > You can also wrap a section of code between method calls of a cProfile.Profile object:

```
profiler = cProfile.Profile()
profiler.enable()
fun()
profiler.disable()
profiler.print_stats(sort='time')
```

128

cProfile

```
[x] [ ] [ ] IPython: chapter1/codes
(hyperf) → codes ipython
Python 3.5.2 |Continuum Analytics, Inc.| (default, Jul 2 2016, 17:53:06)
Type "copyright", "credits" or "license" for more information.

IPython 5.1.0 -- An enhanced Interactive Python.
?           -> Introduction and overview of IPython's features.
%quickref -> Quick reference.
help        -> Python's own help system.
object?    -> Details about 'object', use 'object??' for extra details.

In [1]: from simul import benchmark

In [2]: %prun benchmark()
         707 function calls in 1.231 seconds

Ordered by: internal time

ncalls  tottime  percall  cumtime  percall filename:lineno(function)
      1   1.230    1.230    1.230   1.230 simul.py:21(evolve)
      1   0.000    0.000    0.001   0.001 simul.py:118(<listcomp>)
    300   0.000    0.000    0.000   0.000 random.py:342(uniform)
    100   0.000    0.000    0.000   0.000 simul.py:10(__init__)
    300   0.000    0.000    0.000   0.000 {method 'random' of '_random.Random' objects}
      1   0.000    0.000    1.231   1.231 {built-in method builtins.exec}
      1   0.000    0.000    1.231   1.231 <string:>1(<module>)
      1   0.000    0.000    1.231   1.231 simul.py:117(benchmark)
      1   0.000    0.000    0.000   0.000 simul.py:18(__init__)
      1   0.000    0.000    0.000   0.000 {method 'disable' of '_lsprof.Profiler' objects}

In [3]:
```

129

Profiling memory usage with memory_profiler

```
[x] [ ] [ ] IPython: chapter1/codes
IPython 5.1.0 -- An enhanced Interactive Python.
?           -> Introduction and overview of IPython's features.
%quickref -> Quick reference.
help        -> Python's own help system.
object?    -> Details about 'object', use 'object??' for extra details.

In [1]: %load_ext memory_profiler

In [2]: from simul import benchmark_memory

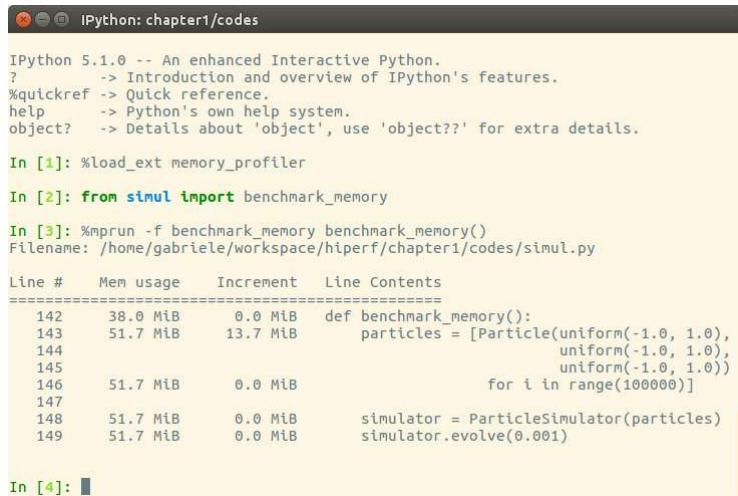
In [3]: %mprun -f benchmark_memory benchmark_memory()
Filename: /home/gabriele/workspace/hyperf/chapter1/codes/simul.py

Line #  Mem usage  Increment  Line Contents
=====  ======  ======  ======
 142     37.8 MiB    0.0 MiB  def benchmark_memory():
 143     61.5 MiB   23.7 MiB      particles = [Particle(uniform(-1.0, 1.0),
 144                                         uniform(-1.0, 1.0),
 145                                         uniform(-1.0, 1.0))
 146     61.5 MiB    0.0 MiB          for i in range(100000)]
 147
 148     61.5 MiB    0.0 MiB      simulator = ParticleSimulator(particles)
 149     61.5 MiB    0.0 MiB      simulator.evolve(0.001)

In [4]:
```

130

Profiling memory usage with memory_profiler



The screenshot shows an IPython notebook interface with the title bar "IPython: chapter1/codes". The notebook contains the following code and output:

```
IPython 5.1.0 -- An enhanced Interactive Python.
?          -> Introduction and overview of IPython's features.
%quickref -> Quick reference.
help       -> Python's own help system.
object?    -> Details about 'object', use 'object??' for extra details.

In [1]: %load_ext memory_profiler

In [2]: from simul import benchmark_memory

In [3]: %mprun -f benchmark_memory benchmark_memory()
Filename: /home/gabriele/workspace/hiperf/chapter1/codes/simul.py

Line #  Mem usage  Increment  Line Contents
===== 
142     38.0 MiB    0.0 MiB  def benchmark_memory():
143      51.7 MiB    13.7 MiB      particles = [Particle(uniform(-1.0, 1.0),
144                                         uniform(-1.0, 1.0),
145                                         uniform(-1.0, 1.0))
146      51.7 MiB    0.0 MiB      for i in range(100000)]
147
148      51.7 MiB    0.0 MiB      simulator = ParticleSimulator(particles)
149      51.7 MiB    0.0 MiB      simulator.evolve(0.001)

In [4]:
```

131

Profiling Metrics

- > **ncalls**
 - Number of calls
- > **tottime**
 - Time spent in a function (excluding subcalls)
- > **percall**
 - Time per call
- > **cumtime**
 - Cumulative time (including subcalls)
- > **filename:lineno(function)**
 - Function location

132

Visualizing Profiling Data

- > Tools for Visualization
 - pstats module
 - SnakeViz
 - PyProf2CallTree
- > Example with SnakeViz:
`pip install snakeviz
snakeviz output.prof`

133

Best Practices for Profiling

- > Profile critical sections of code.
- > Use real-world data inputs.
- > Compare different implementations.
- > Visualize results for deeper insights.
- > Avoid over-optimization.

134

Example: Profiling a Sorting Algorithm

```
import cProfile  
import random  
  
def bubble_sort(arr):  
    n = len(arr)  
    for i in range(n):  
        for j in range(0, n - i - 1):  
            if arr[j] > arr[j + 1]:  
                arr[j], arr[j + 1] = arr[j + 1], arr[j]  
  
data = [random.randint(0, 1000) for _ in range(1000)]  
cProfile.run('bubble_sort(data)')
```

135

Analyzing the Output

5 function calls in 0.050 seconds

Ordered by: standard name

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	0.050	0.050	<string>:1(<module>)
1	0.050	0.050	0.050	0.050	exercise01.py:5(bubble_sort)
1	0.000	0.000	0.050	0.050	{built-in method builtins.exec}
1	0.000	0.000	0.000	0.000	{built-in method builtins.len}
...					

136

Profiling Larger Applications

- > Example: Profiling a Web App

```
import cProfile  
from my_web_app import app  
cProfile.run('app.run()', 'output.prof')
```

- > Visualization

- > Use SnakeViz or KCachegrind

137

Profiling Decorator

```
import cProfile  
import random  
  
def profile(func):  
    def wrapper(*args, **kwargs):  
        profiler = cProfile.Profile()  
        profiler.enable()  
        result = func(*args, **kwargs)  
        profiler.disable()  
        profiler.print_stats(sort='time')  
        return result  
    return wrapper
```

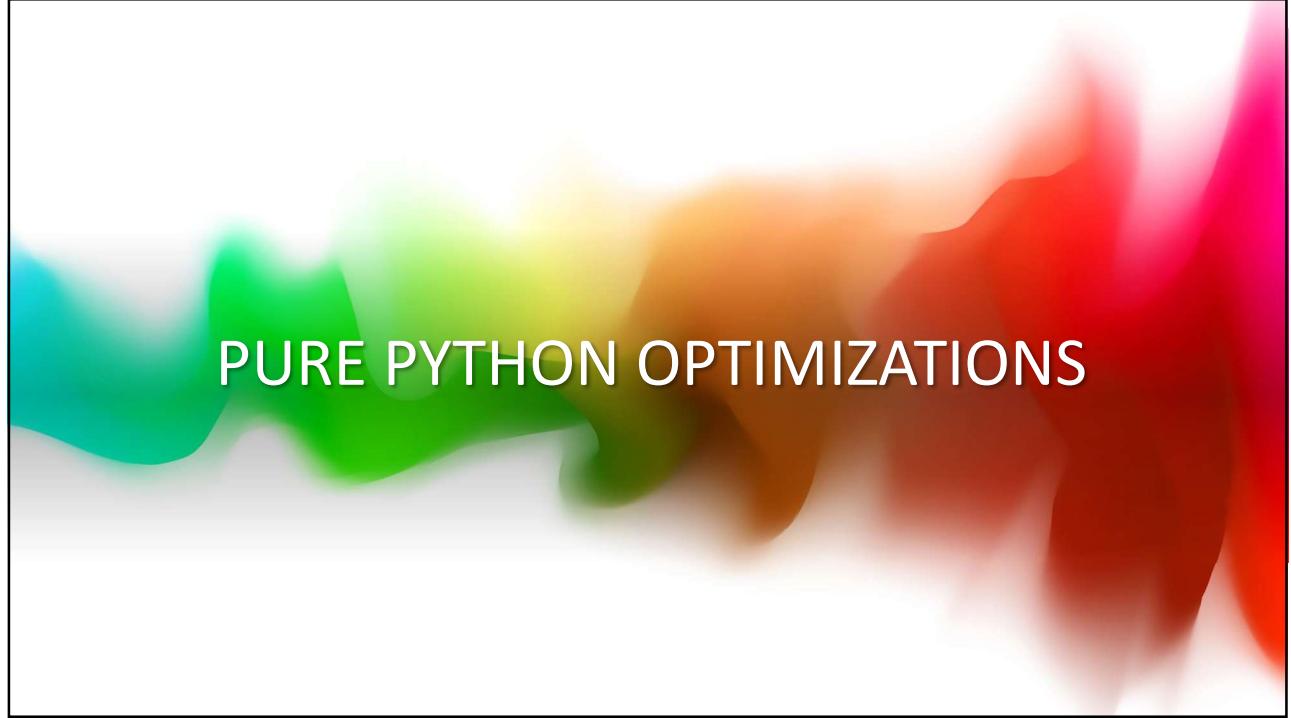
138

Profiling Decorator

```
@profile
def bubble_sort(arr):
    n = len(arr)
    for i in range(n):
        for j in range(0, n - i - 1):
            if arr[j] > arr[j + 1]:
                arr[j], arr[j + 1] = arr[j + 1], arr[j]

data = [random.randint(0, 1000) for _ in range(1000)]
bubble_sort(data)
```

139



PURE PYTHON OPTIMIZATIONS

140

Useful algorithms and data structures

- > Algorithmic improvements are especially effective in increasing performance because they typically allow the application to scale better with increasingly large inputs
- > Efficient Data Structures
 - Choose Wisely
 - Use **set** for membership tests
 - Use **deque** for queue operations
 - Use **defaultdict** for default values

141

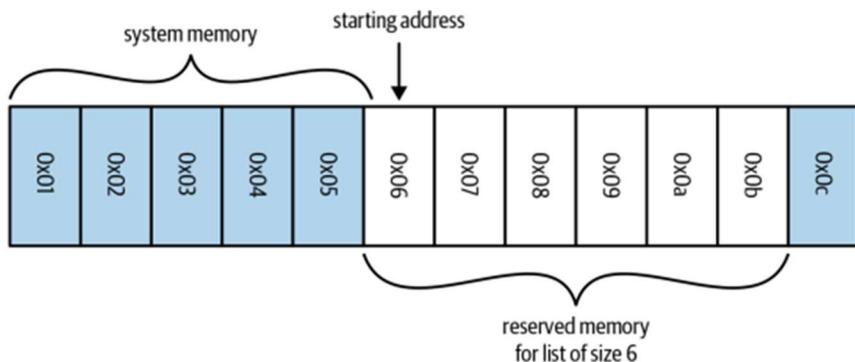
Useful algorithms and data structures

- > Example:
- ```
from collections import defaultdict
counts = defaultdict(int)
for item in data:
 counts[item] += 1
```

142

## Lists and deques

- > Python lists are ordered collections of elements and, in Python, are implemented as resizable arrays.
- > An array is a basic data structure that consists of a series of contiguous memory locations, and each location contains a reference to a Python object.



143

## Lists and deques

```
>>> %%timeit l = list(range(10))
...: l[5]
...
30.1 ns ± 0.996 ns per loop (mean ± std. dev. of 7 runs, 10000000 loops each)

>>> %%timeit l = list(range(10_000_000))
...: l[100_000]
...
28.9 ns ± 0.894 ns per loop (mean ± std. dev. of 7 runs, 10000000 loops each)
```

144

## Linear Search

- > What if we were given an array with an unknown order and wanted to retrieve a particular element?

```
def linear_search(needle, array):
 for i, item in enumerate(array):
 if item == needle:
 return i
 return -1
```

145

## Binary Search

- > Once a list has been sorted, we can find our desired element using a binary search
- > Average case complexity is  $O(\log n)$ .

```
def binary_search(needle, haystack):
 imin, imax = 0, len(haystack)
 while True:
 if imin > imax:
 return -1
 midpoint = (imin + imax) // 2
 if haystack[midpoint] > needle:
 imax = midpoint
 elif haystack[midpoint] < needle:
 imin = midpoint+1
 else:
 return midpoint
```

146

## Lists Versus Tuples

- > If lists and tuples both use the same underlying data structure, what are the differences between the two?
- > Lists are dynamic arrays
  - Mutable and allow for resizing.
- > Tuples are static arrays
  - Immutable
- > Tuples are cached by the Python runtime
  - No need to talk to the kernel to reserve memory every time we want to use one.

147

## Using Built-in Functions

- > Why Built-ins?
  - Implemented in C, faster than Python loops.
- > Examples
  - **sum, min, max, sorted**
  - **map, filter**

148

## Using Built-in Functions

> Example:

```
Inefficient
result = []
for x in data:
 result.append(x**2)
Optimized
result = map(lambda x: x**2, data)
```

149

## List Comprehensions and Generators

- > List Comprehensions
  - Faster and more concise than loops
- > Generators
  - Use less memory by yielding items lazily

150

## List Comprehensions and Generators

> Example:

```
List comprehensions
squares = [x**2 for x in range(10)]
Generator
squares_gen = (x**2 for x in range(10))
```

151

## String Manipulations

> Avoid Concatenation in Loops

– Use **join** for better performance

> Example

```
Inefficient
result = ''
for s in strings:
 result += s
Optimized
result = ''.join(strings)
```

152

## Lazy Evaluation

- > What is Lazy Evaluation?
  - Postponing computations until needed
- > Techniques
  - Use iterators
  - Leverage **itertools** for combinatorics

153

## Lazy Evaluation

- > Example:

```
import itertools
result = itertools.islice(itertools.count(), 10)
```

154

## Exception Handling Optimization

- > Best Practices
  - Avoid using exceptions for flow control
  - Catch specific exceptions.

155

## Exception Handling Optimization

```
> Example:
Inefficient
try:
 result = data[index]
except IndexError:
 result = None
Optimized
result = data[index] if index < len(data) else None
```

156

## Avoiding Global Variables

- > Why Avoid Globals?
  - Slower access compared to local variables.
- > Solution
  - Pass variables as function arguments

157

## Using `@lru_cache`

- > Cache function results to avoid recomputation.

158

## Using @lru\_cache

> Example:

```
from functools import lru_cache
@lru_cache(maxsize=None)
def fibonacci(n):
 if n < 2:
 return n
 return fibonacci(n-1) + fibonacci(n-2)
```

159

## Parallel Processing with multiprocessing

> Utilize multiple CPU cores.

160

## Parallel Processing with multiprocessing

> Example:

```
from multiprocessing import Pool

def square(x):
 return x**2

with Pool(4) as p:
 result = p.map(square, range(10))
```

161

## Using **numpy** for Numerical Computations

> Highly optimized for array operations

162

## Using **numpy** for Numerical Computations

> Example:

```
import numpy as np
data = np.array([1, 2, 3, 4])
result = data * 2
```

163

## Reducing Function Call Overhead

> Inline critical functions to avoid call overhead

164

## Reducing Function Call Overhead

> Example:

```
Inefficient
def add(a, b):
 return a + b

result = add(x, y)

Optimized
result = x + y
```

165

## Optimize Imports

- > Import only what you need
- > Use local imports in specific functions

166

## Avoid Repeated Calculations

```
Inefficient
for i in range(len(data)):
 for j in range(len(data)):
 process(data[i], data[j])

Optimized
length = len(data)
for i in range(length):
 for j in range(length):
 process(data[i], data[j])
```

167

## Avoid Excessive Logging

- > Why?
  - Reduces I/O overhead
- > Best Practices
  - Log only essential information

168

## Use **with** Statement for File Handling

- > Benefits
  - Ensures proper resource cleanup

169

## Use **with** Statement for File Handling

```
Inefficient
f = open('file.txt')
data = f.read()
f.close()

Optimized
with open('file.txt') as f:
 data = f.read()
```

170

## Leverage Type Hints

- > Why?
  - Improves readability and debugging
- > Example

```
def add(a: int, b: int) -> int:
 return a + b
```

171

## Batch Processing

- > Processes large datasets in chunks, reducing processing time and resource usage
- > Decreases overhead caused by multiple I/O operations, such as database writes or network calls.
- > Handles larger datasets by dividing them into manageable pieces, ensuring system stability.
- > Allows the system to focus on bulk operations, reducing the latency associated with individual processing.

172

## When to Use Batch Processing:

- > Database operations
  - Bulk insert or update records instead of performing one at a time.
- > File handling
  - Process multiple files or records together
- > Machine learning
  - Train models on batched data instead of single examples

173

## Batch Database Operations

- > Instead of inserting rows one at a time into a database:

```
Inefficient: Single-row inserts for row in data:
cursor.execute("INSERT INTO table_name VALUES (?, ?)", row)
Optimized: Batch processing
cursor.executemany("INSERT INTO table_name VALUES (?, ?)", data)
connection.commit()
```

174

## Batch File Processing

> Processing a large file line-by-line versus in batches:

```
Inefficient: Line-by-line processing
with open('large_file.txt') as f:
 for line in f:
 process(line)
```

175

## Batch File Processing

> Processing a large file line-by-line versus in batches:

```
Optimized: Batch processing using chunks
def process_in_batches(file_path, batch_size=1000):
 with open(file_path) as f:
 batch = []
 for line in f:
 batch.append(line)
 if len(batch) == batch_size:
 process(batch)
 batch = []
 if batch: # Process remaining lines
 process(batch)

process_in_batches('large_file.txt')
```

176